

# Modeling the exploitation and mitigation of memory safety vulnerabilities

Breakpoint 2012

Matt Miller

Microsoft Security Engineering Center

## Acknowledgements

Richard Tuffin, Julien Vanegue

For their insights and collaboration on classifying memory safety and exploitability

## Related work

Automated/semi-automated exploit generation [[1,2,3,16](#)]

Thanassis Avgerinos, David Brumley, Sean Heelan, Edward Schwartz

Memory safety classification & abstract models for exploitation [[9,10,11,12,13](#)]

Patroklos Argyroudis, Sandeep Bhatkar, Eep Bhatkar, Sergey Bratus, Daniel C. Duvarney, Halvar Flake, Michael E. Locasto, Chariton Karamitas, Meredith L. Patterson, Len Sassaman, R. Sekar, Anna Shubina, Ryan Smith, Chris Valasek

Automated post-mortem analysis of exploitability [[4,15](#)]

Adel Abouchaev, Richard van Eeden, Nitin Kumar Goel, Damian Hasse, Scott Lambert, Lars Opstad, Andy Renk, Jason Shirk, Dave Weinstein, Mark Wodrich, Greg Wroblewski

# Let's start with three assertions

1. Assessing the exploitability of memory safety vulnerabilities is a hard problem

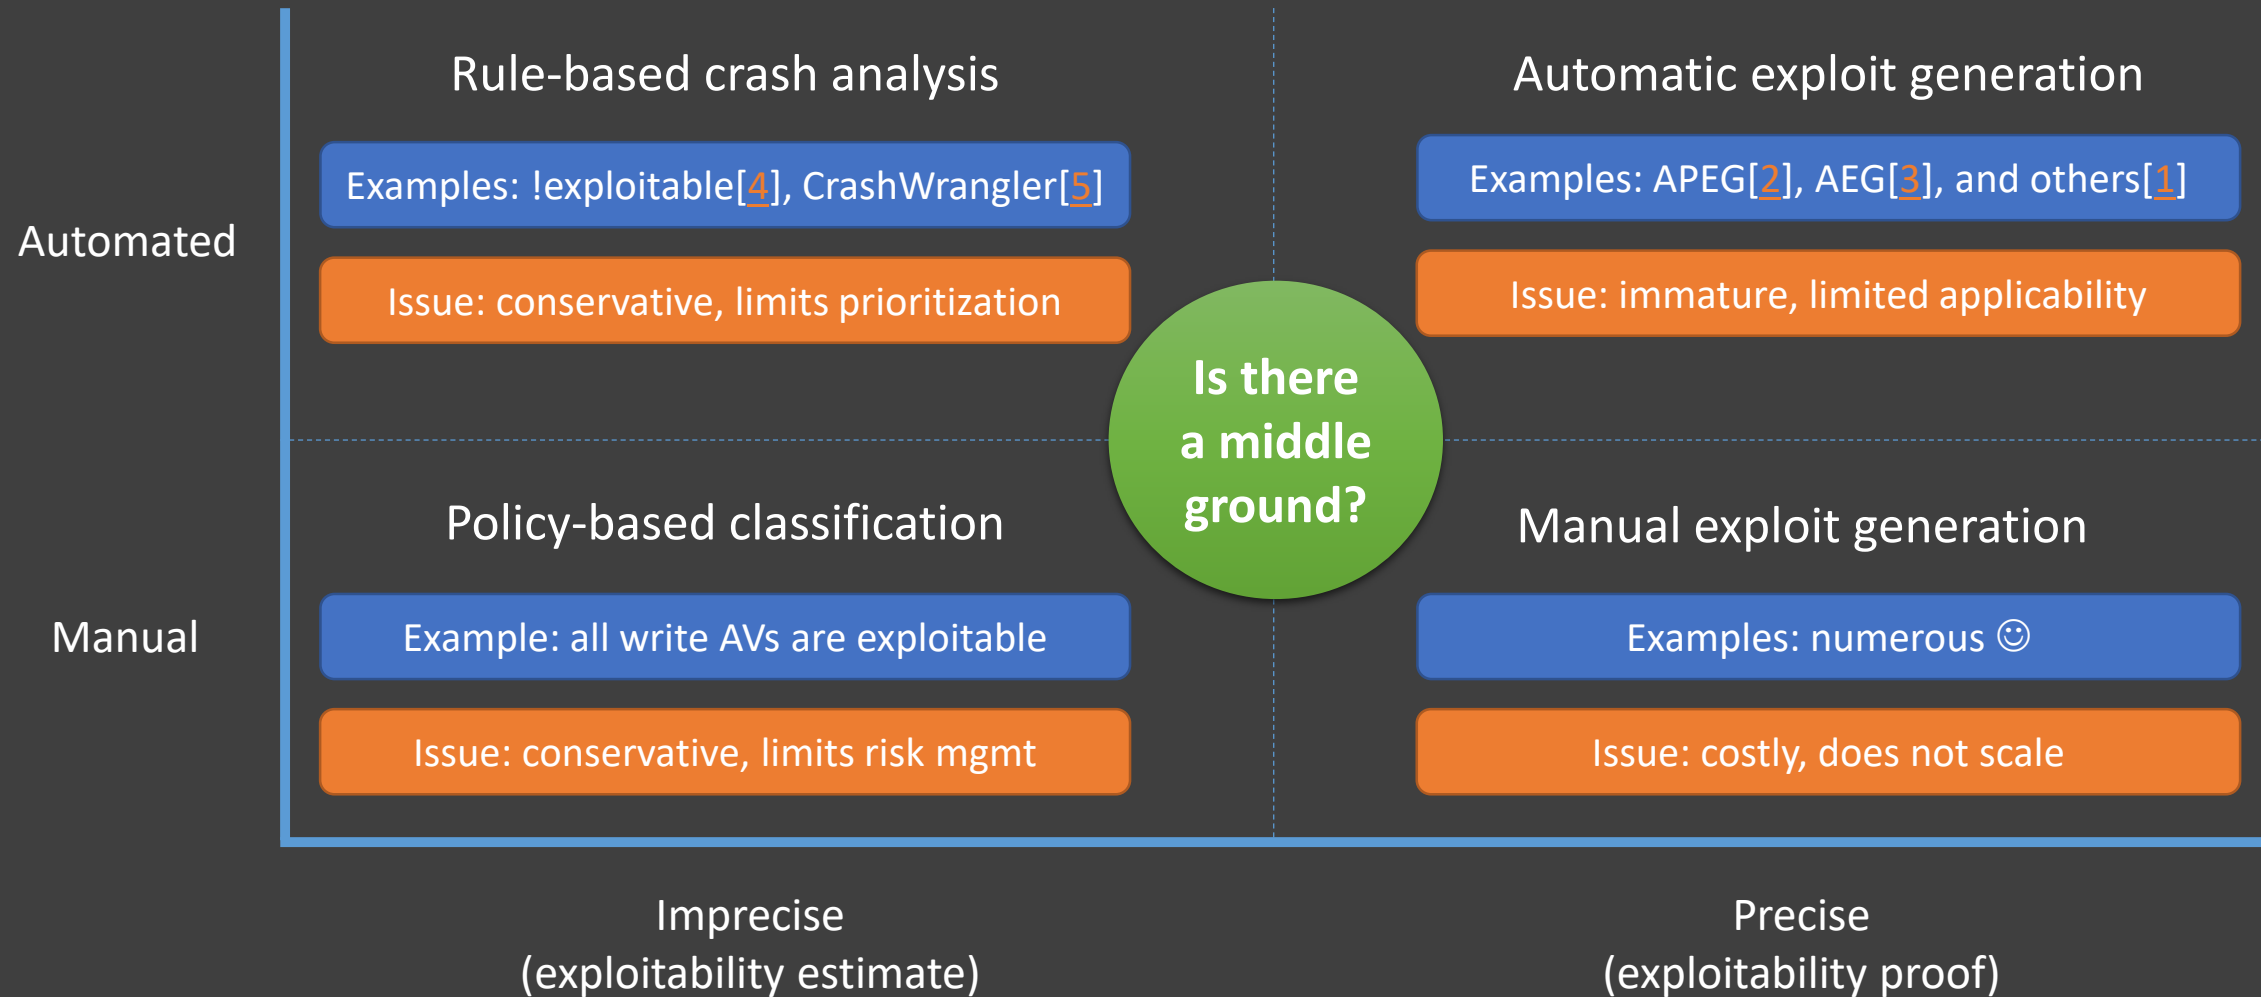
Today, an exploit must be written to prove exploitability (does not scale) or a conservative guess must be made (may overestimate)

2. There is no robust taxonomy for classifying the invariants of a memory safety flaw

CWE[Z] can be used to classify a flaw's type, but not its invariants (which influence the flaw's exploitability)

3. You cannot improve on #1 without first addressing #2

# Current methods for determining exploitability



# Challenges in exploitability assessment

Software vendors must generally be conservative and coarse-grained when estimating exploitability

(!exploitable and Microsoft's Exploitability Index [\[4,6\]](#))

- Risk is often over-estimated[\[14\]](#)
  - Vendor must assume exploitable by default
  - Difficulty of exploitation must be generalized
- Reliant on manual estimation which is non-ideal
  - May be error prone, inconsistent, and hard to verify
  - Individuals must have exploit expertise or be conservative
- Hard to measure impact of mitigations (ex: DEP/ASLR)
- Limits patch deployment prioritization for vendor and customers

Researchers must agree with the vendor's assessment or concretely demonstrate exploitability

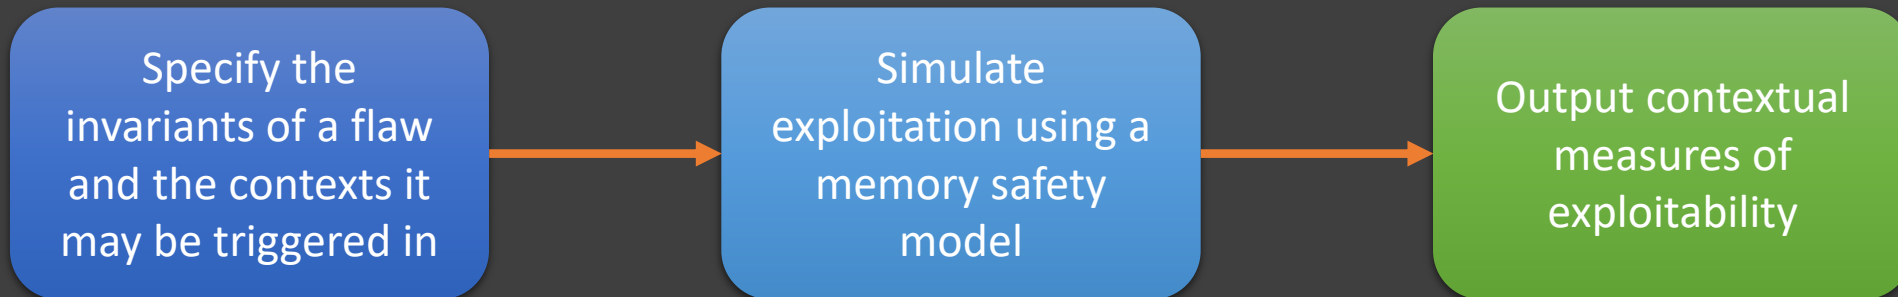
- May create tension between both parties
- May discourage reporting of legitimate issues

# Proposing a way forward

To overcome these challenges, we need a method of assessing exploitability that is...

Accurate	The assessment is correct
Precise	Provides a granular assessment of exploitability that is aware of contextual factors
Objective	The assessment is consistent, repeatable, and can be independently reviewed
Scalable	Human involvement is minimized

An abstract model is one method that could be used to achieve this



# Modeling memory safety

# What is memory safety?

We have many terms for memory safety issues

Buffer overrun, uninitialized use, type confusion, ... [7]

Read AV, write AV, write4, writeN, execute AV, ...

Temporal/spatial memory access errors[8]



We have many terms for memory safety defenses

/GS, /SAFESEH, ...

DEP, ASLR, SEHOP, ...



But how do they map to exploitability?

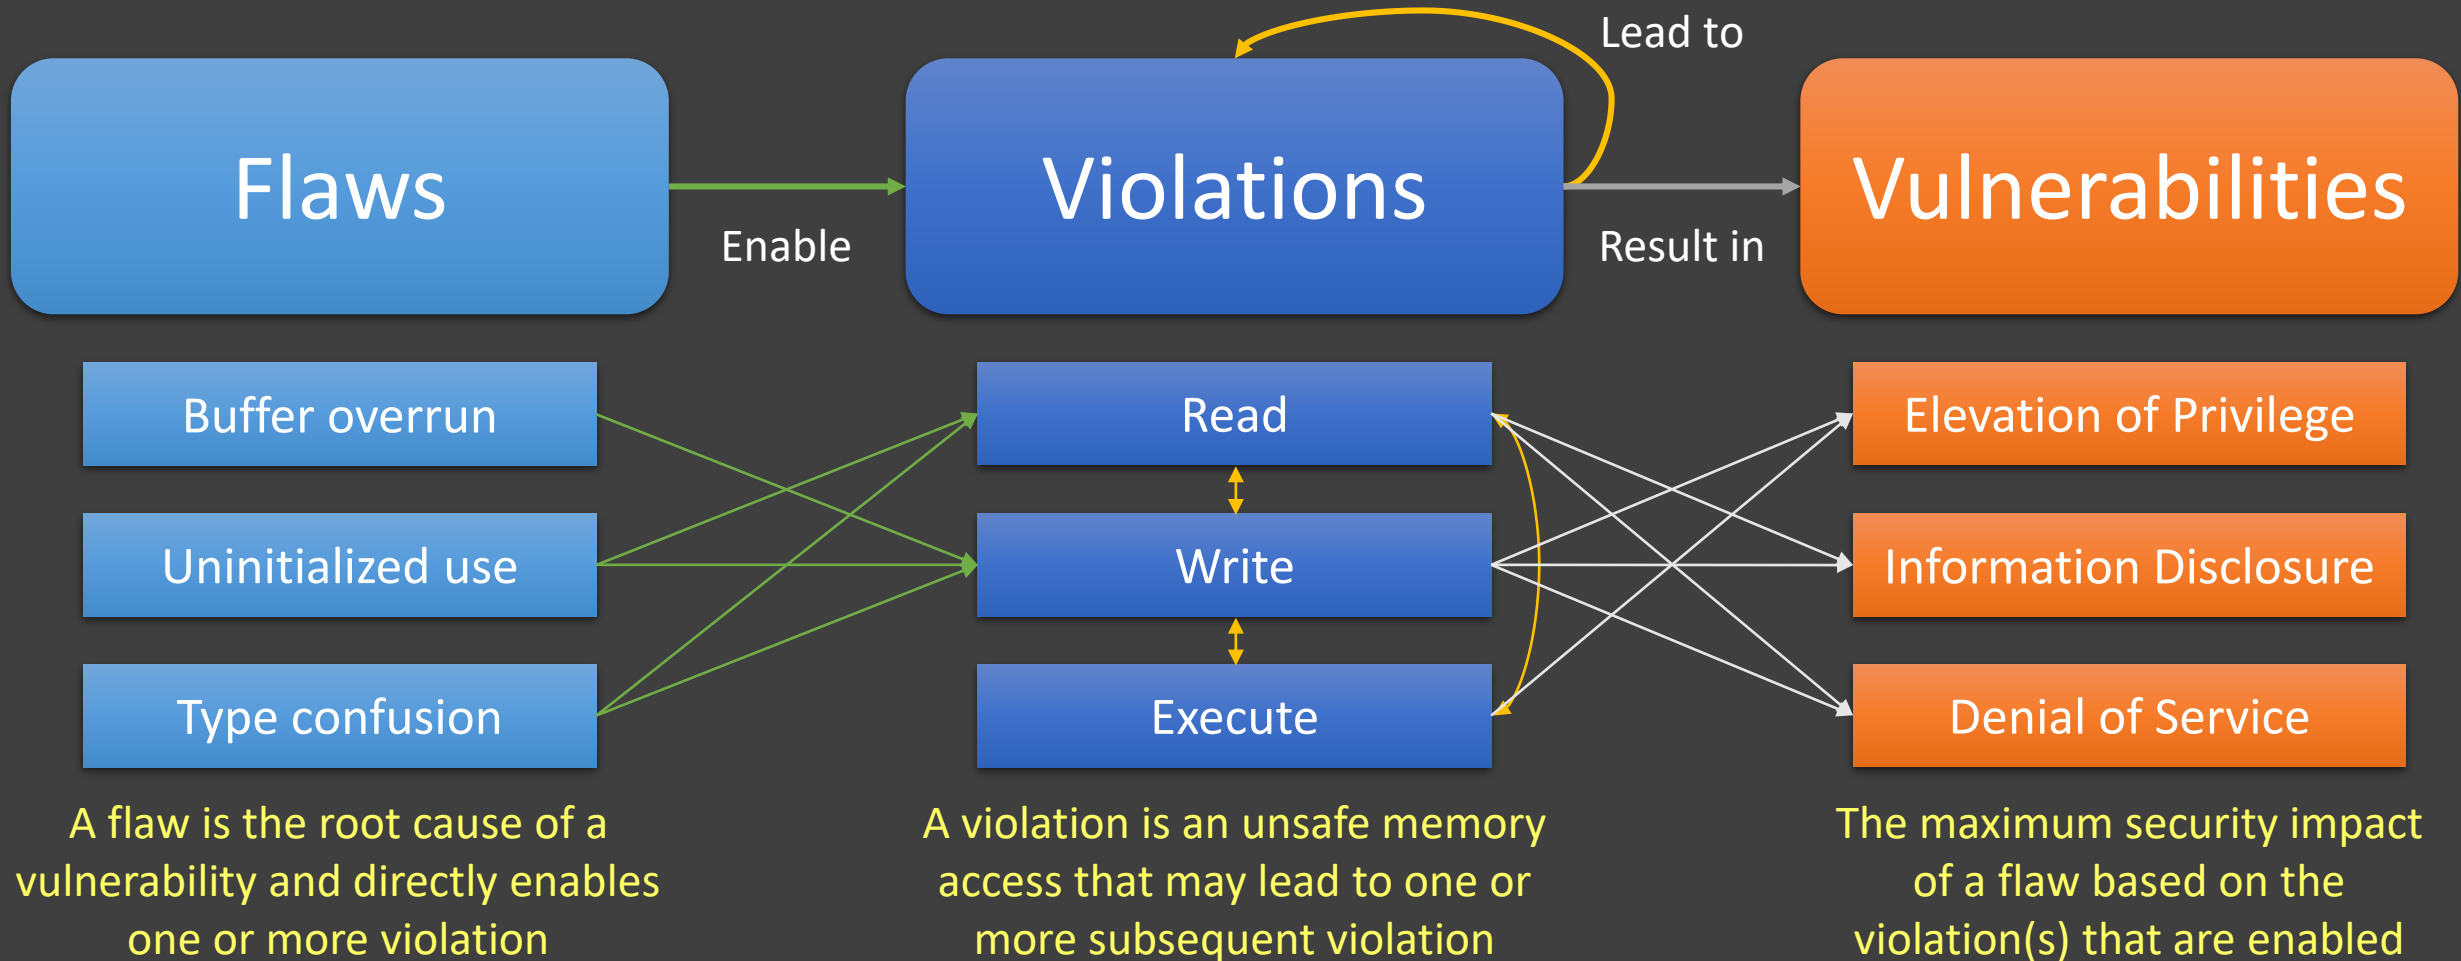


We need a more rigorous taxonomy if we want to be able to measure exploitability



# Fundamental concepts in memory safety

**Memory-safe:** a memory access that is within bounds and refers to an object that is in a valid state



# Classifying memory safety flaws

The taxonomy is generally agreed upon for memory safety flaws (e.g. CWE[7,8])

Arithmetic errors  
and other 2<sup>nd</sup> order  
issues can expose  
memory safety flaws



Boundary error	Type confusion	Uninitialized use
<ul style="list-style-type: none"><li>• Buffer overrun</li><li>• Buffer overread</li><li>• Out-of-bounds array index</li></ul>	<ul style="list-style-type: none"><li>• Invalid type cast</li><li>• Invalid union field access</li></ul>	<ul style="list-style-type: none"><li>• Use after free</li><li>• Double free</li><li>• Uninitialized memory access</li></ul>
Spatial		Temporal



Each flaw can be  
mapped to the set of  
violation(s) that it  
enables

# Classifying memory safety violations

The properties of an unsafe memory access are a convenient way to describe a violation

Two properties establish the basic types of violations

Extended properties add more precision

Memory access method

Memory access parameter state



Parameter	State
Base (b)	Controlled (c)
Content (c)	Fixed (f)
Displacement (d)	Uninitialized (u)
Extent (e)	Unknown (?)

**w-bc-cc-df-ef**

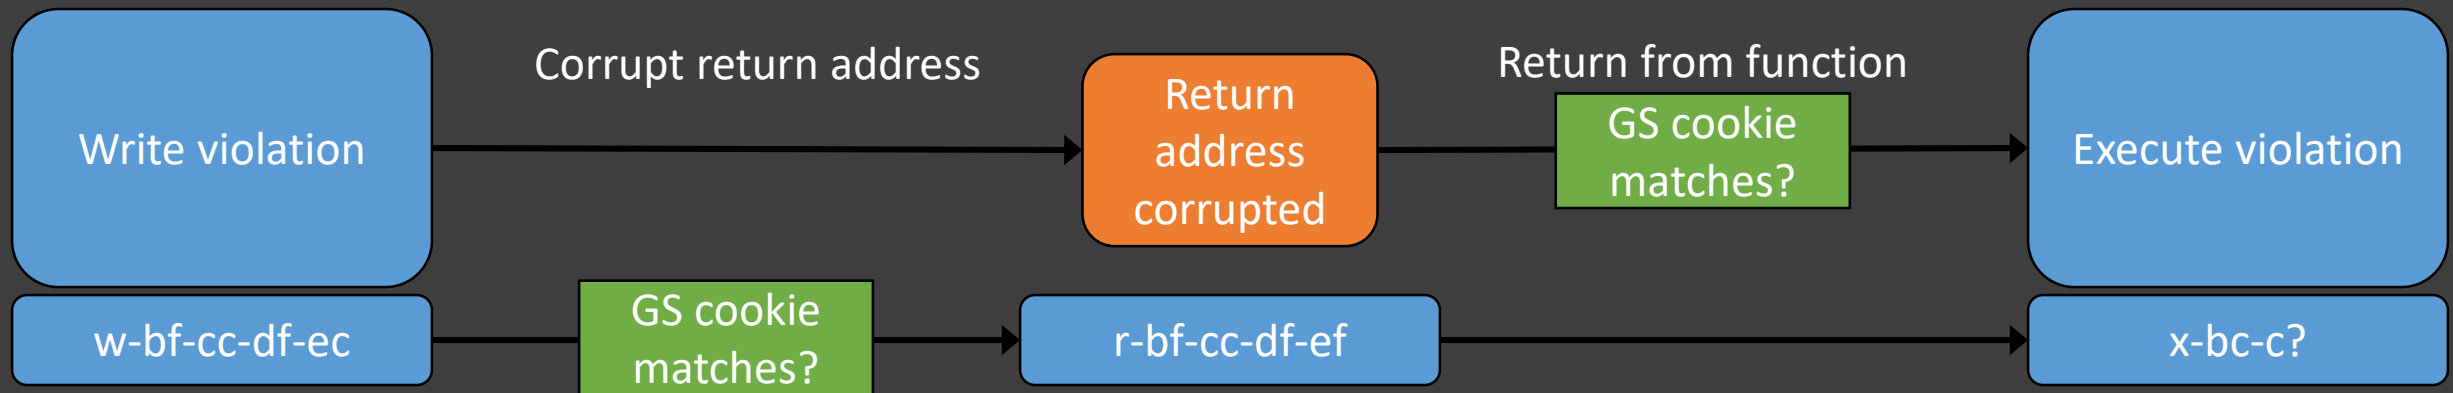
Write, base and content are controlled, displacement and extent are fixed

Property	Values
Addressing mode	absolute, relative
Direction	forward, reverse
Initial displacement	post-adjacent, pre-adjacent, ...
Base region type	heap, stack, ...
Execution domain	user, kernel
Locality	local, remote
...	

# Modeling exploitation & mitigation techniques

An exploitation technique enables a transition from one type of violation to another


Canonical exploitation technique: stack return address overwrite



A mitigation technique introduces new constraints on these transitions

# Exploitation primitives

There is a finite set of primitives for transitioning between the basic types of memory safety violations



Transition	Primitives
$r \rightarrow r$	Read value used as base, displacement, and/or extent of a subsequent read
$r \rightarrow w$	Read value used as base, content, displacement, and/or extent of a subsequent write
$r \rightarrow x$	Read value used as base of an execute
$w \rightarrow r$	Corrupt memory used as base, content, displacement, and/or extent of a subsequent read
$w \rightarrow x$	Corrupt writable code and execute it
$x \rightarrow x$	Execute with a controlled base and/or content

Exploitation techniques combine these primitives in different ways to reach a desired end state

# Vulnerability classification

# A typical vulnerability triage workflow

An uninitialized function ptr is read from the freed object leading to an execute violation. The read does not lead to a subsequent read or write.

Read violation

Execute violation

Write violation

Read violation

Use after free

Write violation

Read violation?

Execute violation?

## Challenges with this workflow

How do you identify the “terrain” of violations to explore?  
How do you explore that terrain consistently and completely?  
How do you deal with and convey unknowns?  
How can you do this without requiring exploitation expertise?

A constant is written to freed mem, but the subsequent violations are unknown as they depend on what is corrupted by the write.

```
Object *p = new Object();  
...  
delete p;  
...  
if (x) { p->value = 0; }  
else { p->OnCompleteCallback(); }
```

1. Determine the root cause (the flaw)

2. Explore violations enabled by the flaw

3. Stop when exploitable violation found

# VEXClass: a vulnerability classification assistant

VEXClass is a proof of concept tool that adds more structure to the vulnerability triage & classification process

Enables uniform and consistent classification

- Well-defined model for flaw and violation invariants
- Supports expression of “known” unknowns
- Classification output is normalized & easily reviewed

Provides a map of the terrain to explore

- Can be used as a “checklist” for the triage process
- Helps ensure that all possibilities are considered
- Analysts do not need state-of-the-art exploitation expertise

Facilitates more precise exploitability analysis

- Classification output is used as input for exploitability analysis
- Separates classification process from exploitability assessment
- Advances in exploitation do not require re-classification





# Example output

```
<VulnerabilityDescription xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance">
  <Flaw>
    <Name>Heap use after free</Name>
    <Symbol>heap_use_after_free</Symbol>
    <Guid>41937a23-9419-4b09-96f7-b0828290a3c4</Guid>
    <Locality>Remote</Locality>
    <AccessRequirement>Unauthenticated</AccessRequirement>
    <ExecutionDomain>User</ExecutionDomain>
    <SourceFile>foo.c</SourceFile>
    <SourceLines>1337</SourceLines>
  </Flaw>
  <Violations>
    <Violation>
      <Name>read uninitialized function ptr</Name>
      <Symbol>r-bf-cu-df-ef</Symbol>
      <Guid>a01d0f76-6b58-4d4d-9e9d-69f1782e9154</Guid>
      <TransitiveViolations>
        <Violation>
          <Name>call through function ptr</Name>
          <Symbol>x-b?-c?</Symbol>
          <Guid>744feefa-0fe4-4ebf-8eba-7610a36cda4f</Guid>
          <TransitiveViolations />
          <Method>Execute</Method>
          <BaseState>Unknown</BaseState>
          <BaseRegionType xsi:nil="true" />
          <ContentState>Unknown</ContentState>
          <ContentDataType xsi:nil="true" />
          <ContentContainerDataType xsi:nil="true" />
          <DestinationContentState>Nonexistent</DestinationContentState>
          <DisplacementState>Nonexistent</DisplacementState>
          <DisplacementInitialOffset xsi:nil="true" />
          <ExtentState>Nonexistent</ExtentState>
          <AddressingMode>Absolute</AddressingMode>
          <Direction xsi:nil="true" />
          <ControlTransferMethod xsi:nil="true" />
        </Violation>
      </TransitiveViolations>
    </Violation>
  </Violations>
</VulnerabilityDescription>
```

Normalized classification  
output enables information  
sharing and exploitability  
analysis

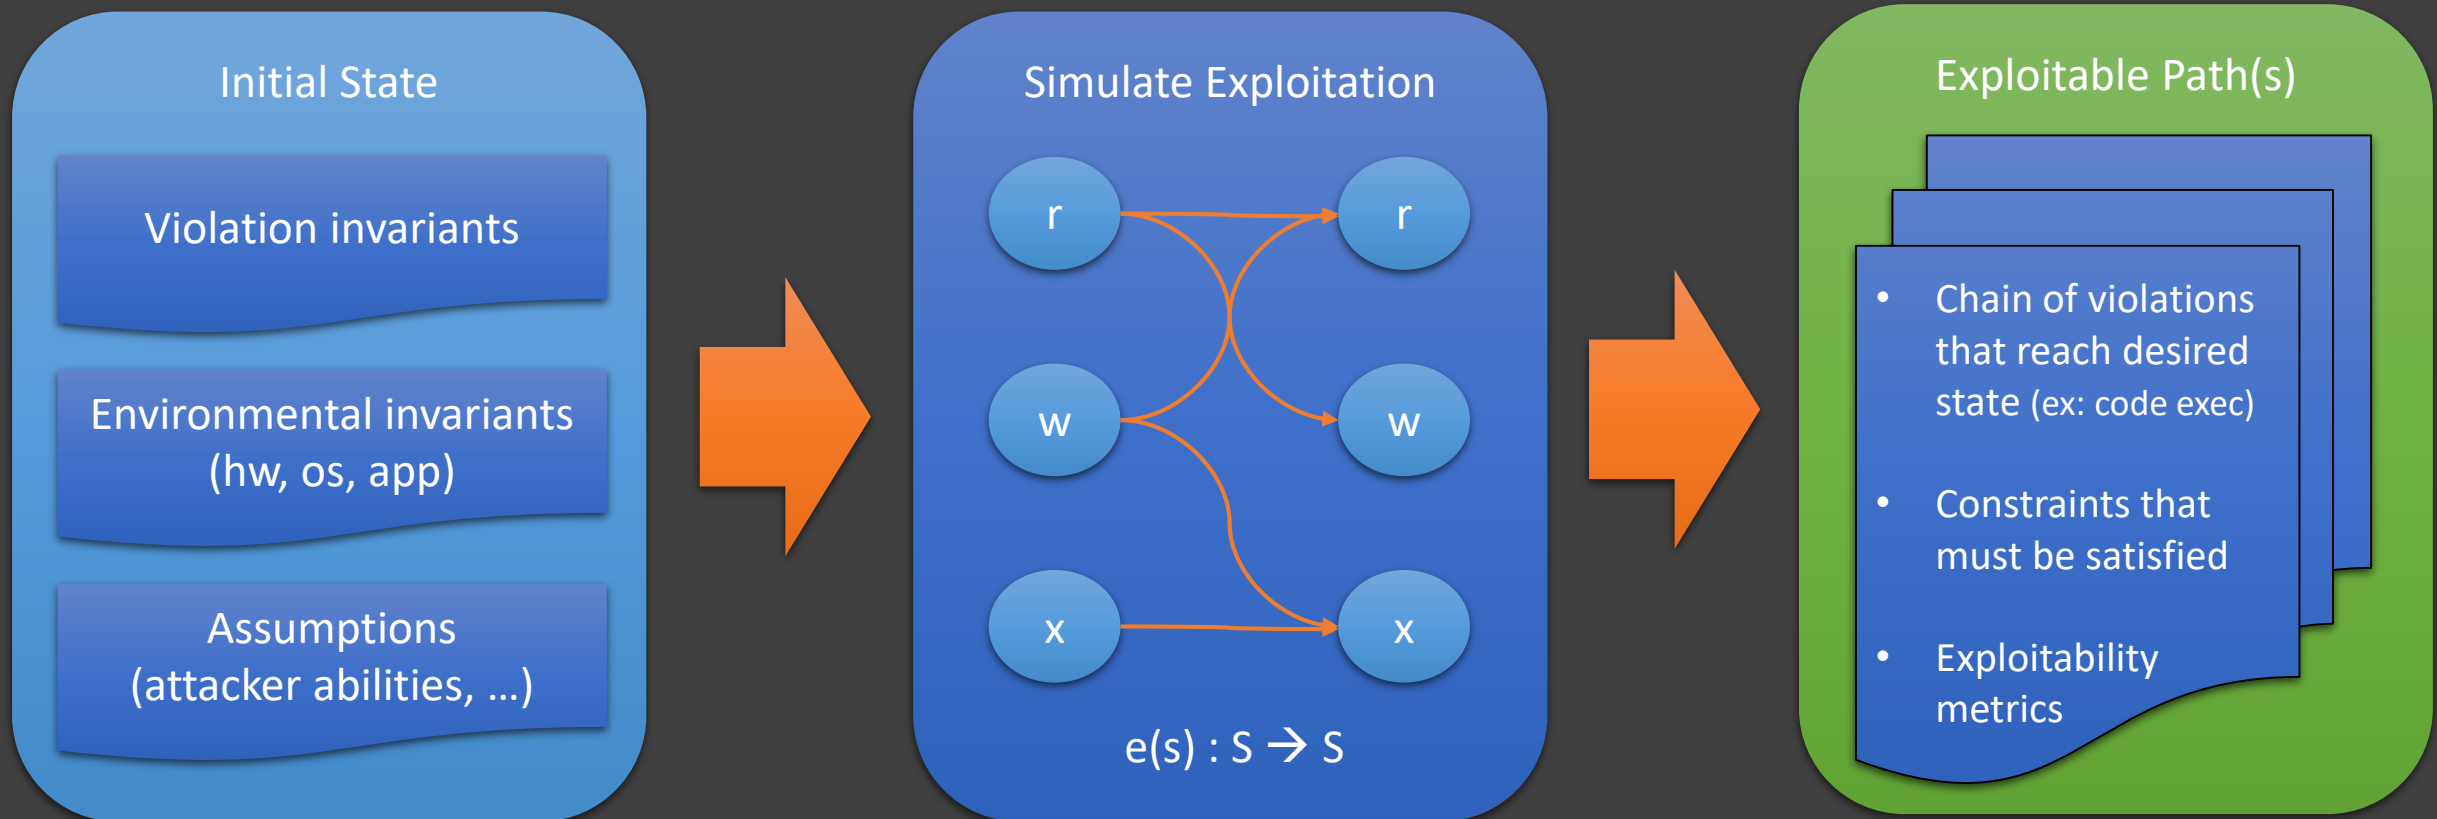
(without giving away  
specific details of the issue)

# VEXClass Demo

# Measuring exploitability

# Measuring exploitability via an abstract model

The invariants of a violation form part of the initial state that is used to measure exploitability



This model quantifies exploitability based on the probability that a sequence of exploitation techniques will succeed


# Simulating exploitation

A state machine (NFA) is a convenient way to model exploitation


Each **state** is composed of the invariants of a violation and the context it is triggered in

Violation	w-bf-cc-df-ec
Application	Internet Explorer 9
OS	Windows 7 SP1
HW	x86 + pae
...	...

Exploitation **techniques** define the transition function from one state to another



Each transition has zero or more constraints that can be **probabilistically** satisfied

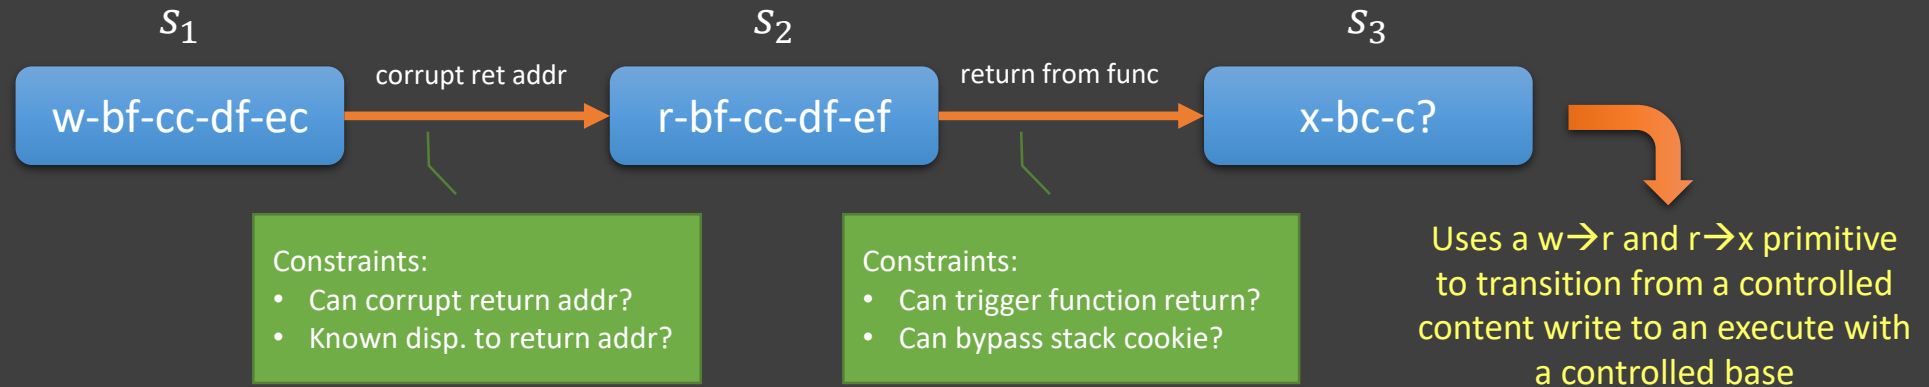


One technique is exploitable 5% of the time, the other not at all

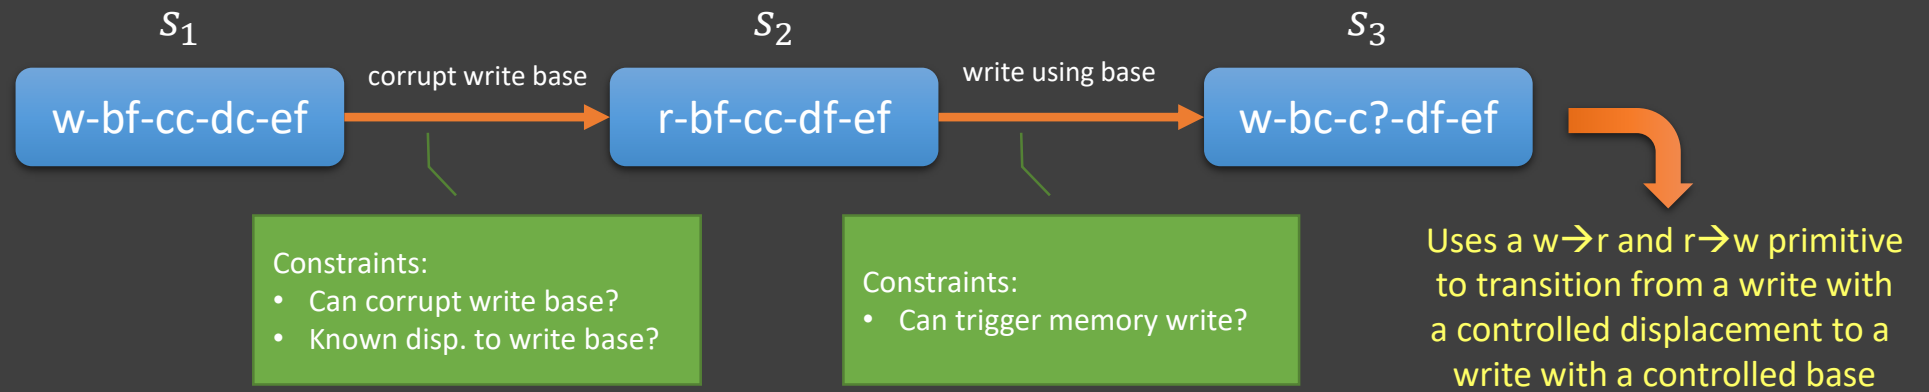
Simulation runs until reaching a fixed point  $[e(s) = s]$  or a desired end state (e.g. code execution)

# Examples of exploitation techniques

Stack return  
address overwrite

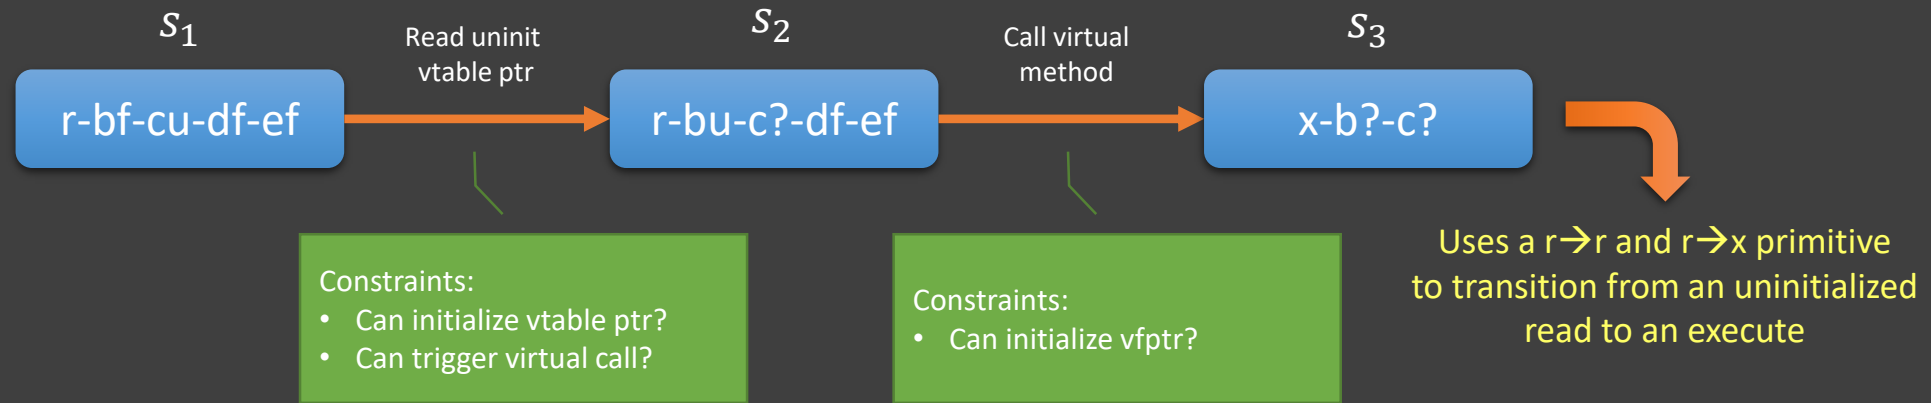


Convert a relative  
write into an  
absolute write

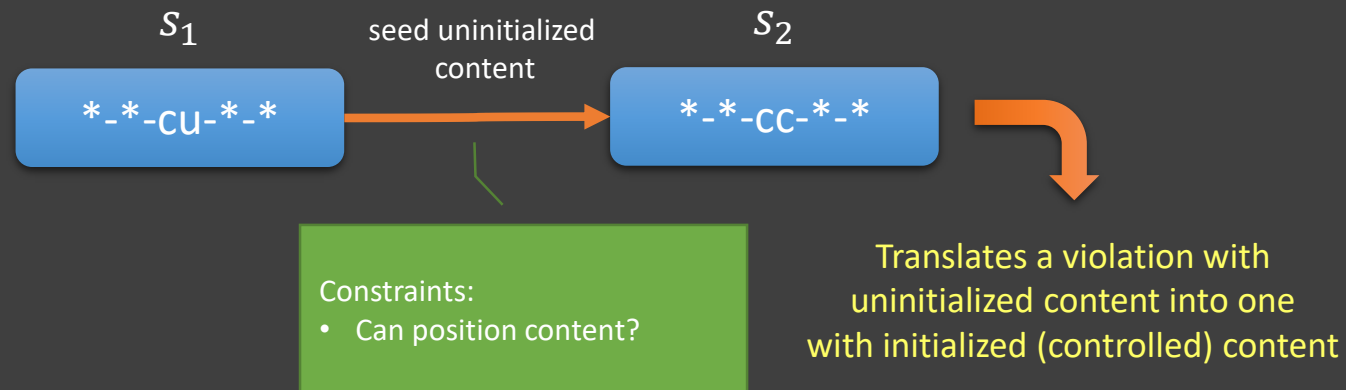


# Examples of exploitation techniques (II)

Corrupt C++ virtual table pointer via use after free



Seed uninitialized content



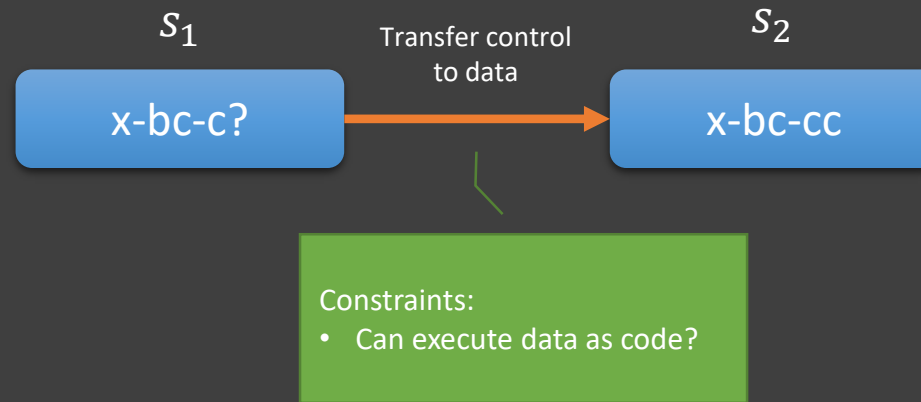
Combined



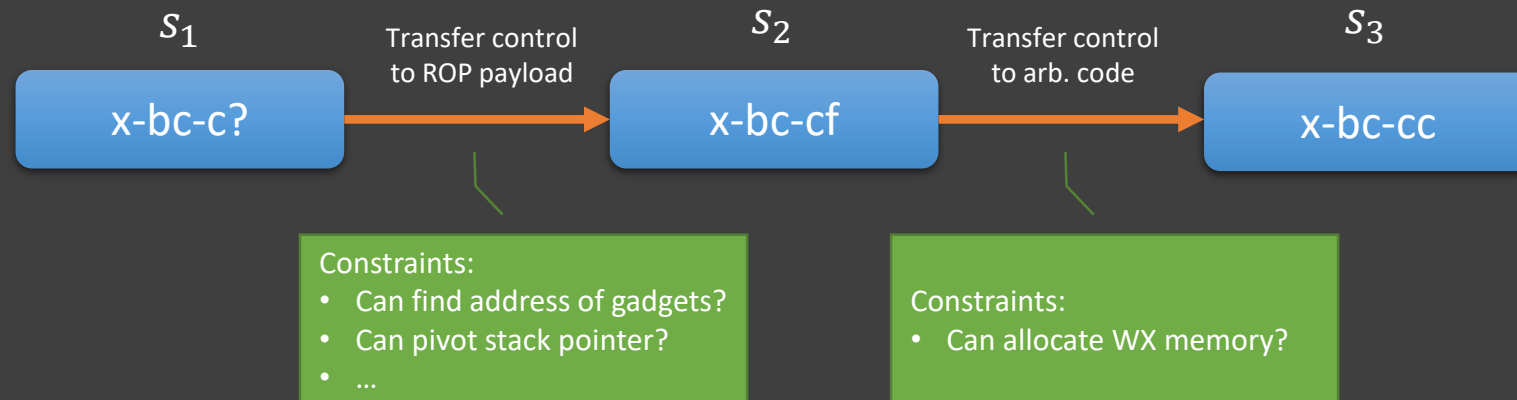


# Examples of exploitation techniques (III)

Execute data as code



Execute ROP stage to arbitrary code



# Chaining exploitation techniques

Exploitation techniques can be chained together to reach a desired end state

## Example #1: stack buffer overrun

From violation	Transition	To violation
w-bf-cc-df-ec	corrupt ret addr	r-b?-cc-d?-e?
r-b?-cc-d?-e?	return from func	x-bc-c?
x-bc-c?	load non-ASLR img	x-bc-cf
x-bc-cf	execute ROP stage	x-bc-cc

## Example #2: C++ object use after free

From violation	Transition	To violation
r-bf-cu-df-ef	init vtable ptr	r-bf-cc-df-ef
r-bf-cc-df-ef	read vtable ptr	r-bc-c?-d?-e?
r-bc-c?-d?-e?	spray vtable	r-bc-cc-d?-e?
r-bc-cc-d?-e?	call virt method	x-bc-c?
x-bc-c?	load non-ASLR img	x-bc-cf
x-bc-cf	execute ROP stage	x-bc-cc

Exploitability of both chains depends on probability of satisfying the constraints of each transition

# Simulation Demo

# Applying this model

Measuring value, guiding investments, and improving risk assessment

# Measuring the value of defensive technologies

No established system exists to measure the impact of defensive technologies like as DEP, ASLR, and /GS

Impact is often measured based on the ability to break techniques used by public exploits[18]

E.g. “exploit X would not have worked with mitigation Y in place”

- Impact is measured based on past behavior
- Absence of relevant exploits can be problematic
  - Hard to justify new tech without data
- Does not quantify difficulty of exploitation *after* enabling mitigation
  - Impact measurement is hindsight only


A model of exploitation can be used to concretely measure the impact of defensive technologies

E.g. “with mitigation Y in place, the exploitability in a given scenario is Z”

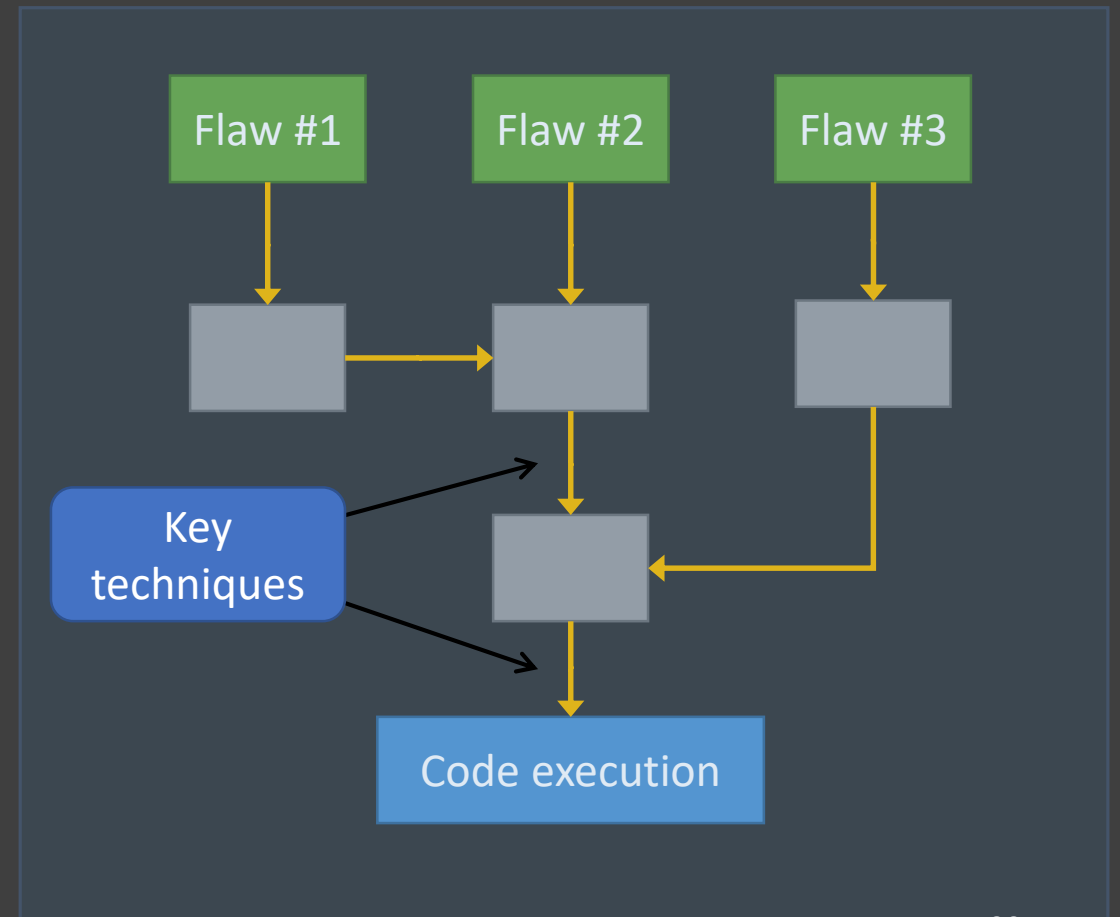
- Impact is measured based on possible behavior
- Open to scrutiny, refinement, and customization
  - Impact may change because of new exploitation techniques
- Facilitates more rigorous “what-if” scenarios
  - E.g. exploring value of a new mitigation

# Guiding defensive investments

Invest in detection & mitigation of flaw types that are most easily exploited & most likely to exist



Invest in mitigations for exploitation techniques that are key to exploiting many types of flaws



# Improving risk management

A model of exploitation could enable more granular and effective risk management

Provide a measure of exploitability for each affected product version rather than generalizing

Enable level of acceptable risk to be customized based on assumptions about an attacker's abilities

Enable continuous and transparent refinement based on changes in the state of the art

Provide an objective description of exploitability that can be more easily agreed upon or refuted

# Conclusion & next steps

## Conclusion

- Measuring risk associated with memory safety vulnerabilities is challenging
  - Often requires conservative analysis which may not reflect actual risk
- An abstract model of memory safety could help improve on this situation
  - Provide a measure of exploitability that enables more effective risk mgmt

## Next steps

- Feedback: this is where you come in 😊
  - Any and all feedback welcome, especially critical feedback!
  - Have something that would be hard to model? I want to hear about it!
- Continue to refine the current prototype of the model described today



Questions?



© 2012 Microsoft Corporation. All rights reserved. Microsoft, Windows, Windows Vista and other product names are or may be registered trademarks and/or trademarks in the U.S. and/or other countries. The information herein is for informational purposes only and represents the current view of Microsoft Corporation as of the date of this presentation. Because Microsoft must respond to changing market conditions, it should not be interpreted to be a commitment on the part of Microsoft, and Microsoft cannot guarantee the accuracy of any information provided after the date of this presentation. MICROSOFT MAKES NO WARRANTIES, EXPRESS, IMPLIED OR STATUTORY, AS TO THE INFORMATION IN THIS PRESENTATION.

# References

- [1] Sean Heelan. *Automatic Generation of Control Flow Hijacking Exploits for Software Vulnerabilities*. PhD thesis. Sep, 2009. <http://seanhn.files.wordpress.com/2009/09/thesis1.pdf>.
- [2] David Brumley, Pongsin Poosankam, Dawn Song, and Jiang Zheng. *Automatic Patch-Based Exploit Generation*. In Proceedings of the IEEE Symposium on Security and Privacy, May 2008. <http://www.cs.cmu.edu/~dbrumley/pubs/apeg.pdf>.
- [3] Thanassis Avgerinos, Sang Kil Cha, Brent Lim Tze Hao, David Brumley. *AEG: Automatic Exploit Generation*. Network and Distributed System Security Symposium, Feb 2011. <http://security.ece.cmu.edu/aeg/aeg-ndss-2011.pdf>.
- [4] Microsoft Security Engineering Center (MSEC). *!exploitable Crash Analyzer*. Jun, 2009. <http://msecdbg.codeplex.com/>.
- [5] Apple Product Security. *Announcing Crashwrangler*. Jul, 2009. <http://seclists.org/dailydave/2009/q3/11>.
- [6] Microsoft. *Microsoft Exploitability Index*. Oct, 2008. <http://technet.microsoft.com/en-us/security/cc998259.aspx>.
- [7] MITRE. *Common Weakness Enumeration (CWE)*. <http://cwe.mitre.org/data/index.html>.
- [8] Sandeep Bhatkar, Eep Bhatkar, R. Sekar, Daniel C. Duvarney. *Efficient Techniques for Comprehensive Protection from Memory Error Exploits*. USENIX Security, 2005. <http://citeseerx.ist.psu.edu/viewdoc/summary?doi=10.1.1.59.9963>.
- [9] Matt Miller. *State of the Exploit*. ToorCon Seattle, 2008. [http://hick.org/~mmiller/presentations/seatoorcon08/2008\\_seatoorcon.ppt](http://hick.org/~mmiller/presentations/seatoorcon08/2008_seatoorcon.ppt).
- [10] Sergey Bratus, Michael E. Locasto, Meredith L. Patterson, Len Sassaman, Anna Shubina. *Exploit Programming*. 2011. <http://www.cs.dartmouth.edu/~sergey/langsec/papers/Bratus.pdf>.
- [11] Chris Valasek and Ryan Smith. *Exploitation in the Modern Era*. Black Hat Europe, 2011.
- [12] Halvar Flake. *Exploitation and state machines*. Infiltrate, 2011. [http://immunityinc.com/infiltrate/archives/Fundamentals\\_of\\_exploitation\\_revisited.pdf](http://immunityinc.com/infiltrate/archives/Fundamentals_of_exploitation_revisited.pdf).
- [13] Patroklos Argyroudis, Chariton Karamitas. *Heap Exploitation Abstraction by Example*. OWSAP AppSecResearch, 2012. <http://census-labs.com/media/heap-owasp-appsec-2012.pdf>.
- [14] Microsoft. *Microsoft Security Intelligence Report: Volume 7*. 2009. [http://download.microsoft.com/download/A/3/0/A30A60D9-1303-4B6A-91B7-BB24E0211B05/Microsoft\\_Security\\_Intelligence\\_Report\\_volume\\_7\\_Jan-Jun2009.pdf](http://download.microsoft.com/download/A/3/0/A30A60D9-1303-4B6A-91B7-BB24E0211B05/Microsoft_Security_Intelligence_Report_volume_7_Jan-Jun2009.pdf).
- [15] Adel Abouchaev, Damian Hasse, Scott Lambert, Greg Wroblewski. *Analyze Crashes to Find Security Vulnerabilities in Your Apps*. MSDN Magazine, Nov, 2007. <http://msdn.microsoft.com/en-us/magazine/cc163311.aspx>.

# References

- [16] Edward J. Schwartz, Thanassis Avgerinos, David Brumley. *Q: Exploit Hardening Made Easy*. In the Proceedings of the 2011 USENIX Security Symposium, August, 2011. <http://www.ece.cmu.edu/~ejschwar/papers/usenix11.pdf>
- [17] Jonathan D. Pincus, Brandon Baker. *Beyond Stack Smashing: Recent Advances in Exploiting Buffer Overruns*. IEEE Security & Privacy, vol 2, 2004. <http://ieeexplore.ieee.org/xpl/articleDetails.jsp?arnumber=1324594>.
- [18] Dan Guido. *Exploit Intelligence*. SOURCE Boston, 2011. <http://www.isecpartners.com/storage/docs/presentations/EIP-final.pdf>.