



Toward mitigating arbitrary native code execution in Windows 10

Matt Miller, Microsoft Security Response Center (MSRC)

@epakskape

May, 2017

This presentation is for informational purposes only. MICROSOFT MAKES NO WARRANTIES, EXPRESS OR IMPLIED, IN THIS SUMMARY.



Agenda

How we think about mitigating software vulnerabilities

Our progress toward mitigating native code execution

Impact, lessons learned, and future plans

Imagine you could change the laws of physics

Vertical engineering is a super power when it comes to defense



Change or enable apps to better protect themselves

Microsoft Azure

Change the cloud platform to protect online services & data



Change the OS to protect apps, data, & devices

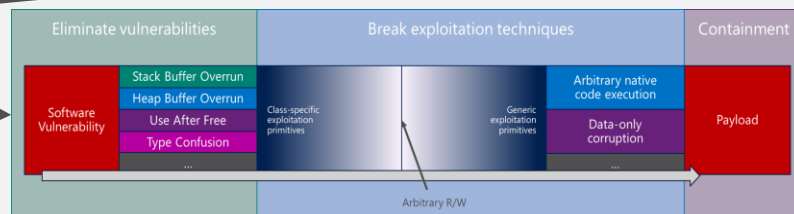
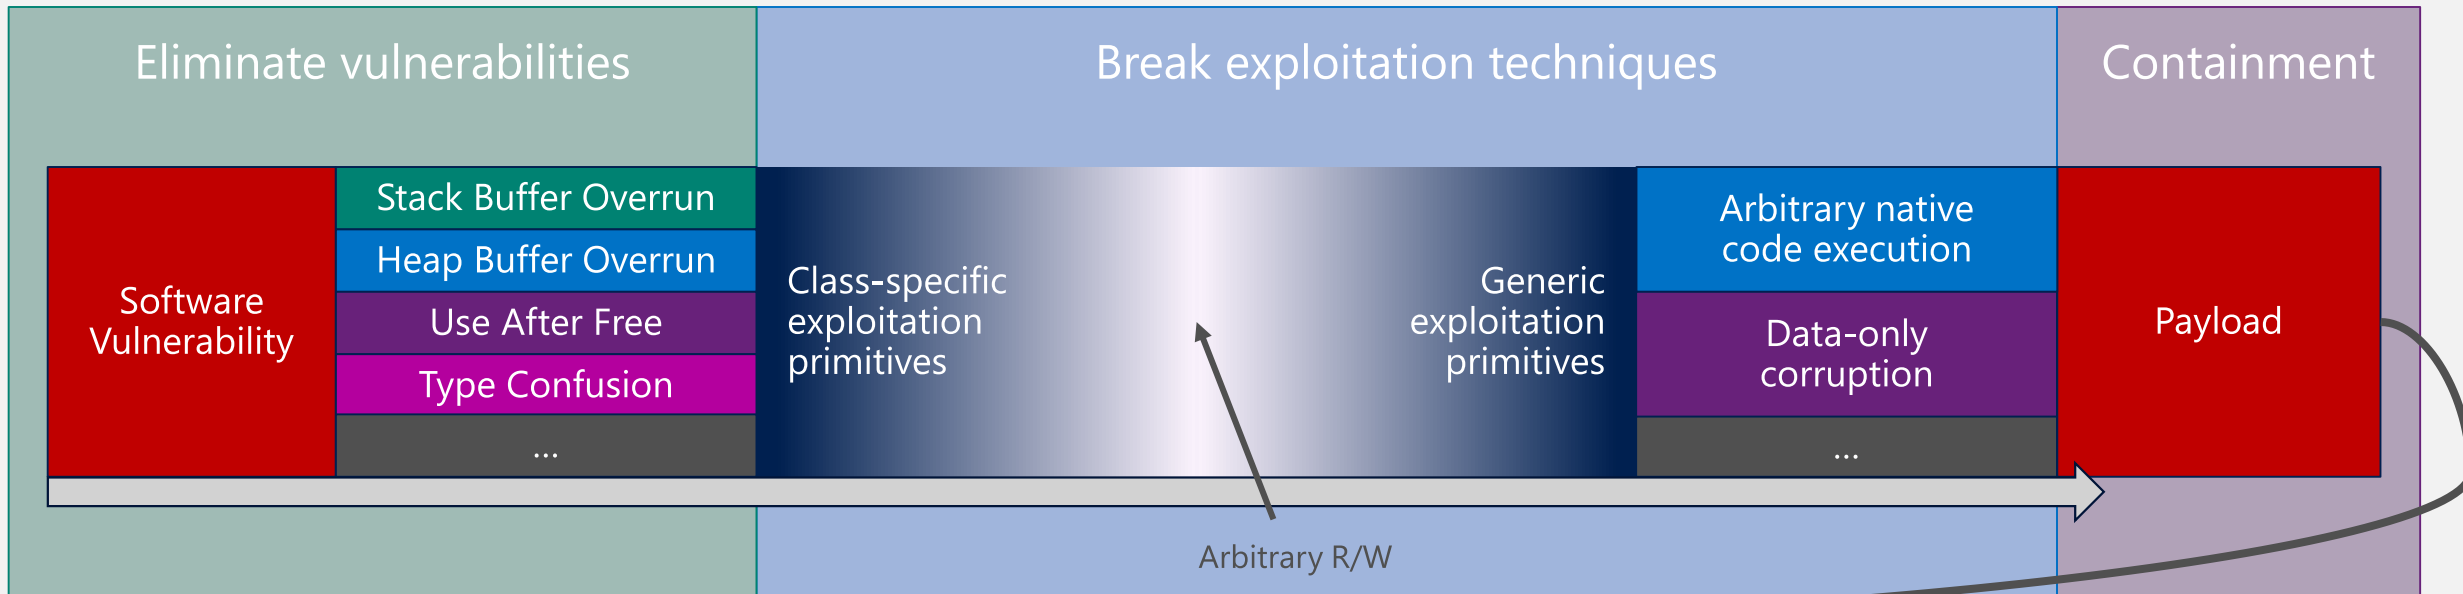


Change the compiler to produce safer code

Collaborate with hardware partners to improve the foundation

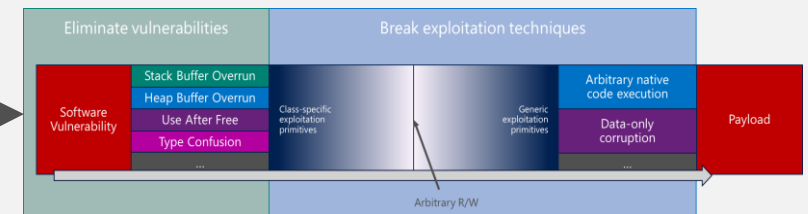
How we think about mitigating software vulnerabilities

Attackers transform software vulnerabilities into tools for delivering a payload to a target device



Attackers typically need to elevate privileges

This means applying the same defenses to privileged attack surfaces



At some point, we lose containment as a defense

This leaves eliminating vulnerabilities & breaking techniques

Layered, data-driven software defense in Windows 10

Our Strategy

Make it difficult & costly to find, exploit, and leverage software vulnerabilities

Our Tactics

Eliminate entire classes of vulnerabilities

Break exploitation techniques

Contain damage & prevent persistence

Limit the window of opportunity to exploit

Today, we'll be focusing on one aspect of breaking exploitation techniques

Chokepoints for breaking exploitation techniques

All exploits rely on combining various primitives to enable the delivery of a payload

Break exploitation techniques

Class-specific
exploitation
primitives

Weak mitigations

Generic
exploitation
primitives

Arbitrary native
code execution

Data-only
corruption

...

Chokepoint #2

Break generic primitives with the assumption that an attacker has arbitrary R/W.

Goal: make it difficult or impossible to deliver a payload independent of the type of vulnerability.

Arbitrary native code execution

- Generic, vulnerability-independent primitive
- Universal & flexible way to execute a payload
- Ubiquitously relied upon by exploits

Good candidate for mitigations 😊

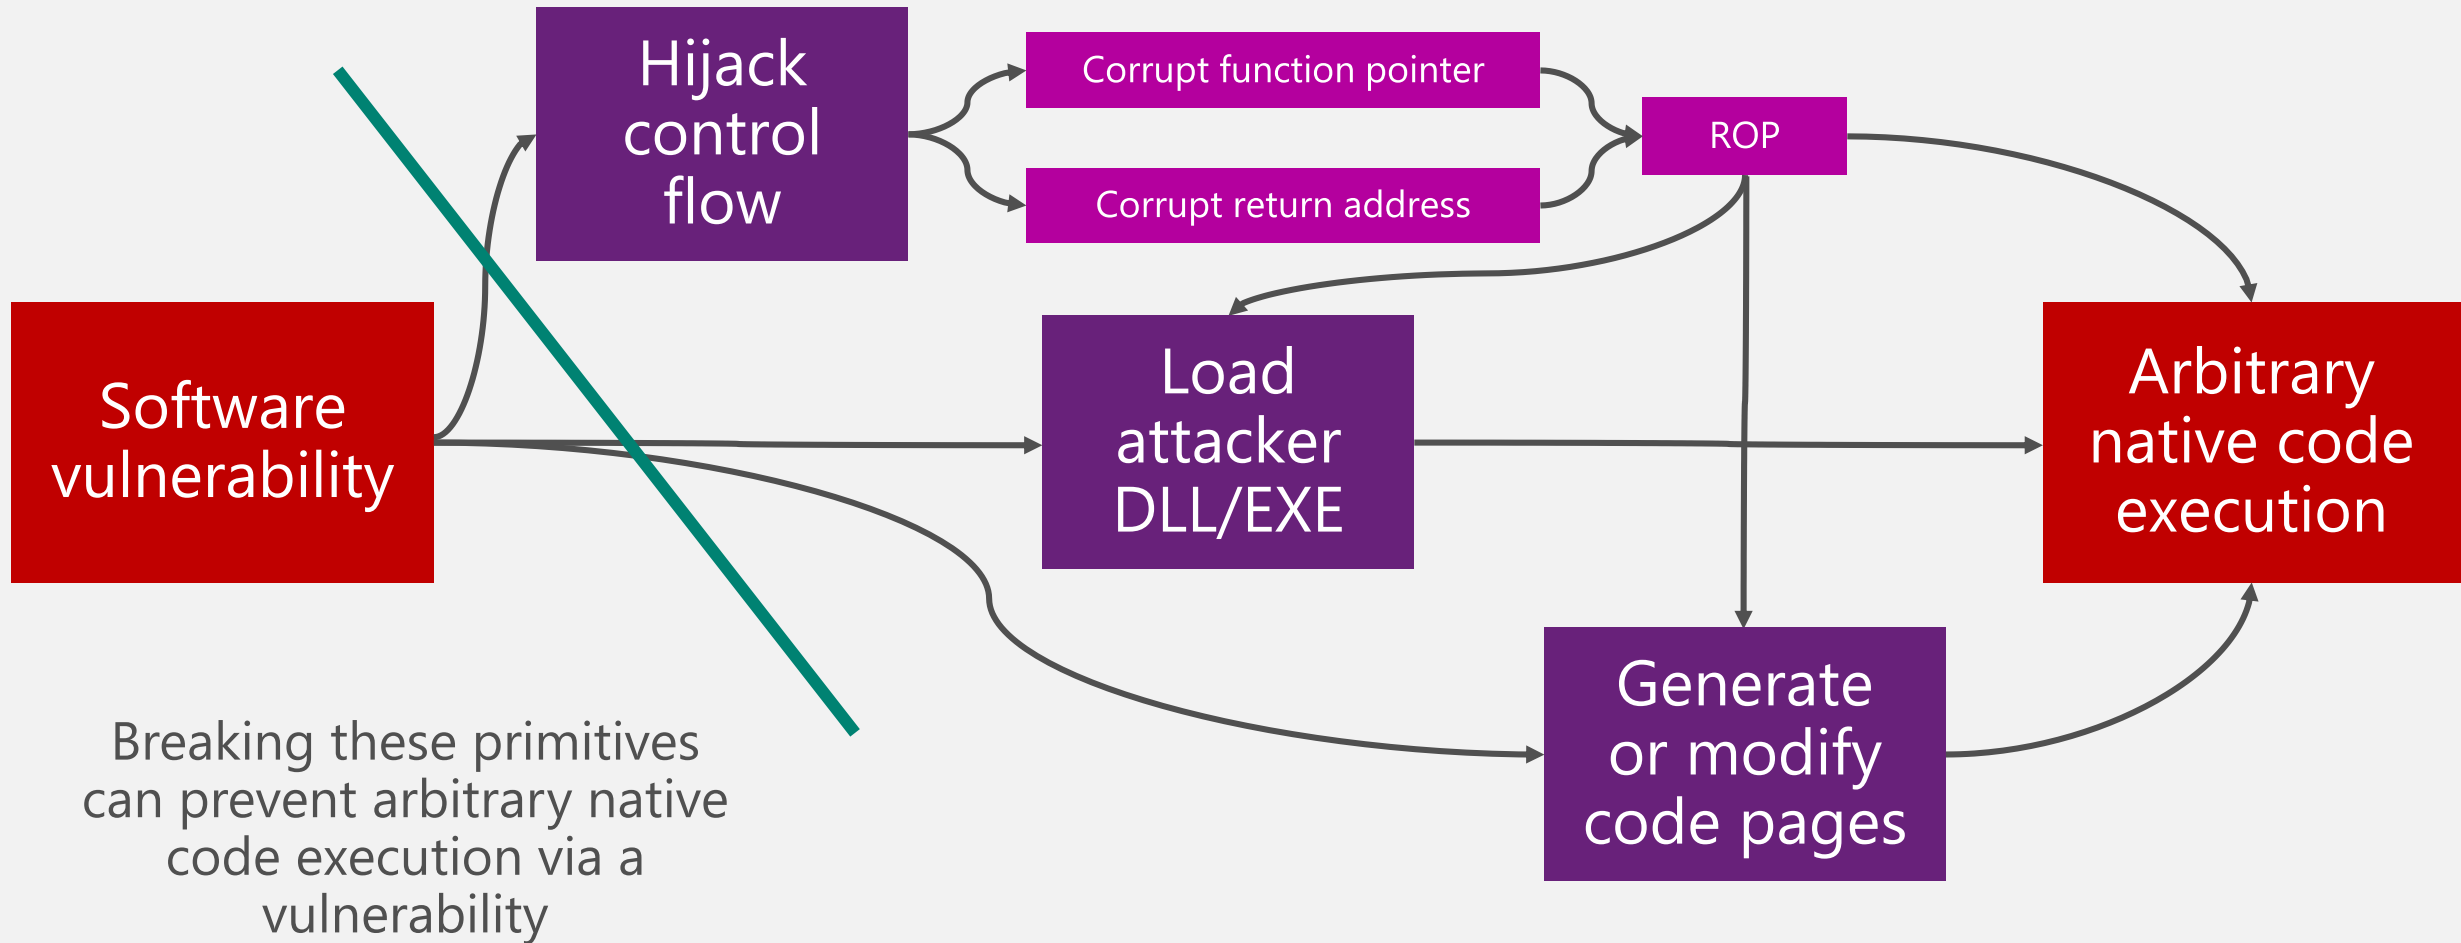
Chokepoint #1

Break class-specific techniques for transforming a vulnerability into generic primitives.

Goal: make it difficult or impossible to exploit certain types of vulnerabilities.

The paths to arbitrary native code execution

There are a finite number of ways to transform a vulnerability into arbitrary native code execution



Mitigating arbitrary native
code execution

Defining our threat model & assumptions

It is critical to define the key aspects of the threat model that we want to defend against

A vulnerability exists

There is a memory corruption vulnerability that could potentially be exploited

Interactive runtime is present

The attacker is able to leverage an interactive runtime (such as a script engine) in conjunction with their exploit

Arbitrary R/W at arbitrary times

The attacker can read/write arbitrary memory locations, with controlled data, at arbitrary points in time

Address space layout is known

The attacker knows where all discoverable memory regions are located – traditional ASLR is a non-factor

This threat model sets a very high bar, but it is grounded by real-world data & expectations

Technologies for mitigating code execution

Prevent
arbitrary code
generation

Code Integrity Guard

Images must be signed and load
from valid places

Arbitrary Code Guard

Prevent dynamic code generation,
modification, and execution

Prevent
control-flow
hijacking

Control Flow Guard

Enforce control flow integrity
on indirect function calls

???

Enforce control flow integrity on
function returns

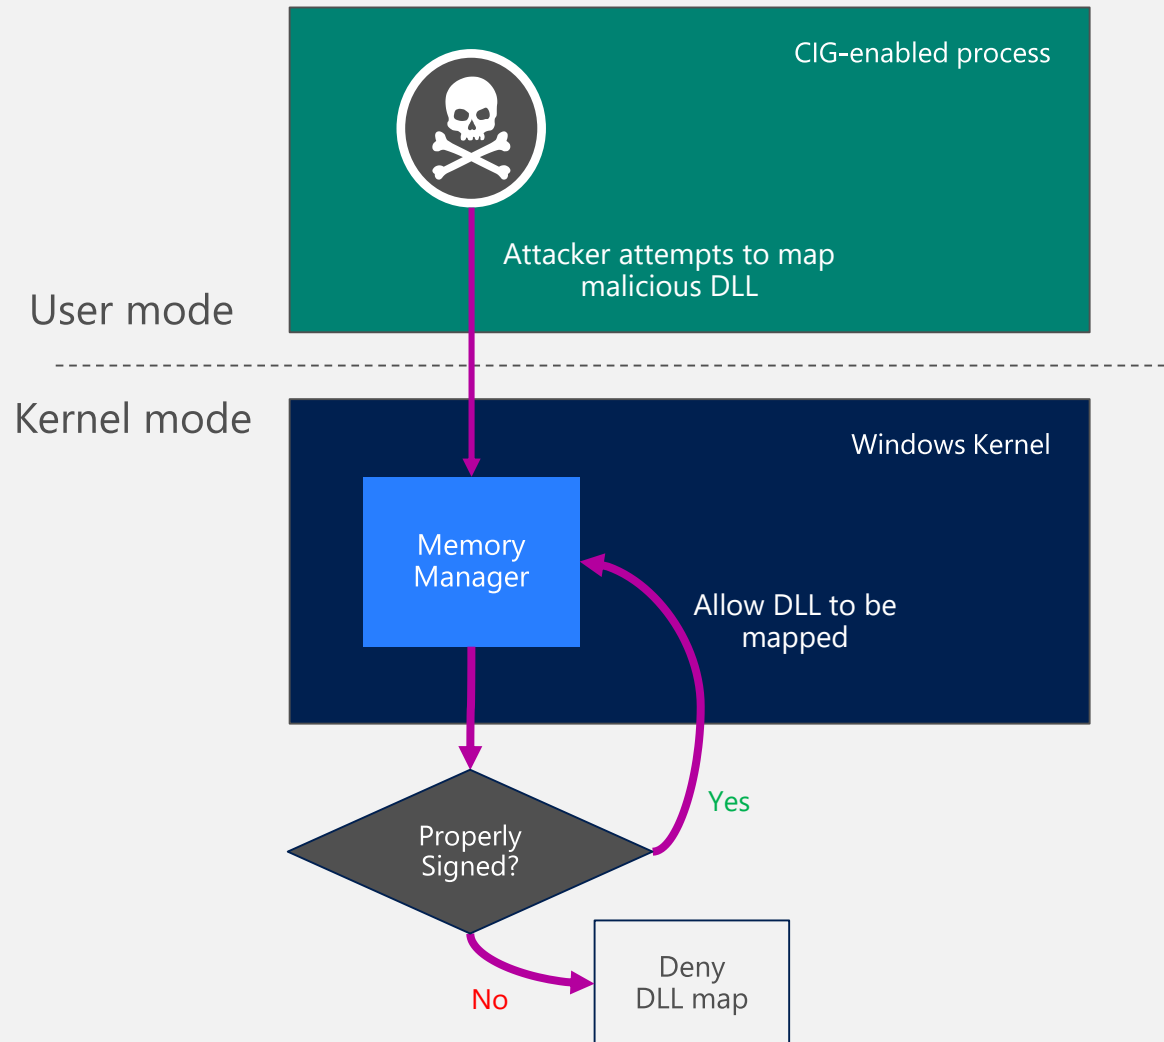
✓ Only valid, signed code pages can
be mapped by the app

✓ Code pages are immutable and
cannot be modified by the app

✓ Code execution stays “on the rails”
per the control-flow integrity policy

Code Integrity Guard (CIG)

CIG leverages code signing restrictions to prevent attacker DLLs from loading



Example of such an attack [provided by Yang Yu](#) @ Black Hat USA 2014

“LoadLibrary” via JavaScript

1. Download a DLL by XMLHttpRequest object, the file will be temporarily saved in the cache directory of IE;
2. Use "Scripting.FileSystemObject" to search the cache directory to find that DLL;
3. Use "Scripting.FileSystemObject" to create a directory named "System32", copy the DLL into that directory, and named it as "shell32.dll";
4. Modify the "SystemRoot" environment variable of current process via "WScript.Shell" object to the upper directory of the "System32" directory created just now;
5. Create "Shell.Application" object, trigger to loading "%SystemRoot%\System32\shell32.dll".



Code Integrity Policies

Only allow Microsoft-signed DLLs

Only allow Microsoft, Store, or WHQL signed DLLs

Configurable via Device Guard policy

Considerations for enabling CIG for an app

3rd party improperly signed binaries that need to be loaded

- Trade-off between security and functionality
- Microsoft Edge will currently disable CIG when 3rd party IMEs are installed

Arbitrary Code Guard (ACG)

Exploits often rely on creating or modifying executable pages to execute their payload

ACG enables two kernel-enforced W^X policies

- ✓ Code is immutable
- ✓ Data cannot become code



Code Integrity Guard (CIg)

- ✓ Only properly signed code pages can be mapped

The following will fail with `ERROR_DYNAMIC_CODE_BLOCKED`

```
VirtualProtect(codePage, ..., PAGE_EXECUTE_READWRITE)
```

```
VirtualProtect(codePage, ..., PAGE_READWRITE)
```

```
VirtualAlloc(..., PAGE_EXECUTE*)
```

```
VirtualProtect(dataPage, ..., PAGE_EXECUTE*)
```

```
MapViewOfFile(hPagefileSection, FILE_MAP_EXECUTE, ...)
```

```
WriteProcessMemory(codePage, ...)
```

```
...
```

ACG + CIg ensure
the integrity of code
pages

Considerations for enabling ACG for an app

**Just-in-Time (JIT)
compilation no longer
works in-proc**

- JITs need to move out-of-process or be disabled
- Microsoft Edge has moved JITs out-of-process in Windows 10 1703

**3rd party binaries may be
incompatible**

- Binaries must enable /DYNAMICBASE, cannot have RWX sections, and must not merge their import table into a code section
- BinSkim has checks for this, see <https://github.com/Microsoft/binskim>
- Microsoft Edge will currently disable ACG when incompatible graphics drivers may be present

**Does not prevent code
injection from privileged
processes**

- Privileged code running outside of the app can still inject dynamic code into an ACG-enabled process

Control Flow Guard (CFG)

Exploits have typically relied on hijacking control-flow through an indirect call

Example control-flow hijack via indirect call to a ROP gadget[1]

```
/* Corrupt sound object vtable ptr */
while (1)
{
    if (this.s[index][j] == 0x00010c00 && this.s[index][j+0x09] == 0x1234)
    {
        soundobjref = this.s[index][j+0x0A];
        dec = soundobjref-cvaddr-1;
        this.s[index][dec/4-2] = cvaddr+2*4+4*4;
        break;
    }

    j++;
}
```

```
/* Run Payload */
this.sound.toString();
```

Transfers control to a
stack pivot ROP gadget



With CFG in place, traditional ROP gadgets and other invalid functions cannot be called indirectly

CFG implements **coarse-grained control-flow integrity** for indirect calls

Compile time

```
void Foo(...) {
    // SomeFunc is address-taken
    // and may be called indirectly
    Object->FuncPtr = SomeFunc;
}
```

Metadata is automatically added to the image which identifies functions that may be called indirectly

```
void Bar(...) {
    // Compiler-inserted check to
    // verify call target is valid
    _guard_check_icall(Object->FuncPtr);
    Object->FuncPtr(xyz);
}
```

A lightweight check is inserted prior to indirect calls which will verify that the call target is valid at runtime

Runtime

Image Load

- Update valid call target data with metadata from PE image

Process Start

- Map valid call target data

Indirect Call

- Perform O(1) validity check
- Terminate process if invalid target

[1] <https://github.com/rapid7/metasploit-framework/blob/abd76c50000e75bcac0616b96cd8583e1df3927f/external/source/exploits/CVE-2014-0322/AsXploit.as>

Control Flow Guard Bypasses & Enhancements

Like all mitigations, CFG has by design limitations that place constraints on its overall effectiveness

- ✓ Return addresses are not protected
- ✓ Valid functions can be called out of context
- ✓ “Fail-open” design for compatibility

Since shipping CFG, researchers have identified bypasses and additional enhancements have been made

Bypass	Status
Non-enlightened Just-in-Time (JIT) compilers	Mitigated in latest version of Edge on Windows 10 (Chakra, Adobe Flash, and WARP)
Multiple non-instrumented indirect calls reported to our Mitigation Bypass Bounty	Mitigated in latest version of Edge on Windows 10
Calling sensitive APIs out of context	NtContinue/longjmp – mitigated for all CFG enabled apps on Windows 10
	VirtualProtect/VirtualAlloc – mitigated in latest version of Microsoft Edge on Windows 10
	LoadLibrary – mitigated in latest version of Microsoft Edge on Windows 10 via CIG
	WinExec – mitigated in Edge on Windows 10 anniversary edition via child process policy
	All exports – mitigated in Edge on Windows 10 1703 via export suppression
Corrupting mutable read-only memory	Known limitation that we are exploring solutions for

What about return addresses?

With CFG in place, we have observed attackers shifting to target return addresses

Protecting return addresses is a decades old (and very difficult) problem

- ✓ Stack cookies (/GS, etc)
 - Only mitigates stack buffer overruns
- ✓ Shadow/split stack solutions
 - E.g. SafeStack (<http://clang.llvm.org/docs/SafeStack.html>)
- ✓ Fine-grained CFI solutions
 - E.g. RAP (https://www.grsecurity.net/rap_announce.php)

Why have none of the generic solutions achieved mainstream adoption yet?

Because software engineering is hard ☺

Protecting return addresses on Windows means satisfying many engineering requirements

Security	Must be robust against our threat model
Performance	Cost must be within reason and proportional to value
Compatibility	Don't break legacy apps
Interoperability	Binary compatibility with previous <i>and</i> future versions of Windows
ABI compliant	We can't rebuild the world
Agility	Don't paint ourselves into a corner
Developer friction	Cost for developers to enable should be minimal (ideally zero)

...and satisfying all of these requirements is *very* hard ☺

Return address protection

We have worked with Intel on designing a hardware-assisted solution for return address protection

Our red team also came up with a clever solution that works on existing x86_64 hardware

- ✓ Control-flow Enforcement Technology (CET)
 - Indirect branch tracking via ENDBRANCH
 - Return address protection via a shadow stack
- ✓ Hardware-assists for helping to mitigate control-flow hijacking & ROP
- ✓ Robust against our threat model

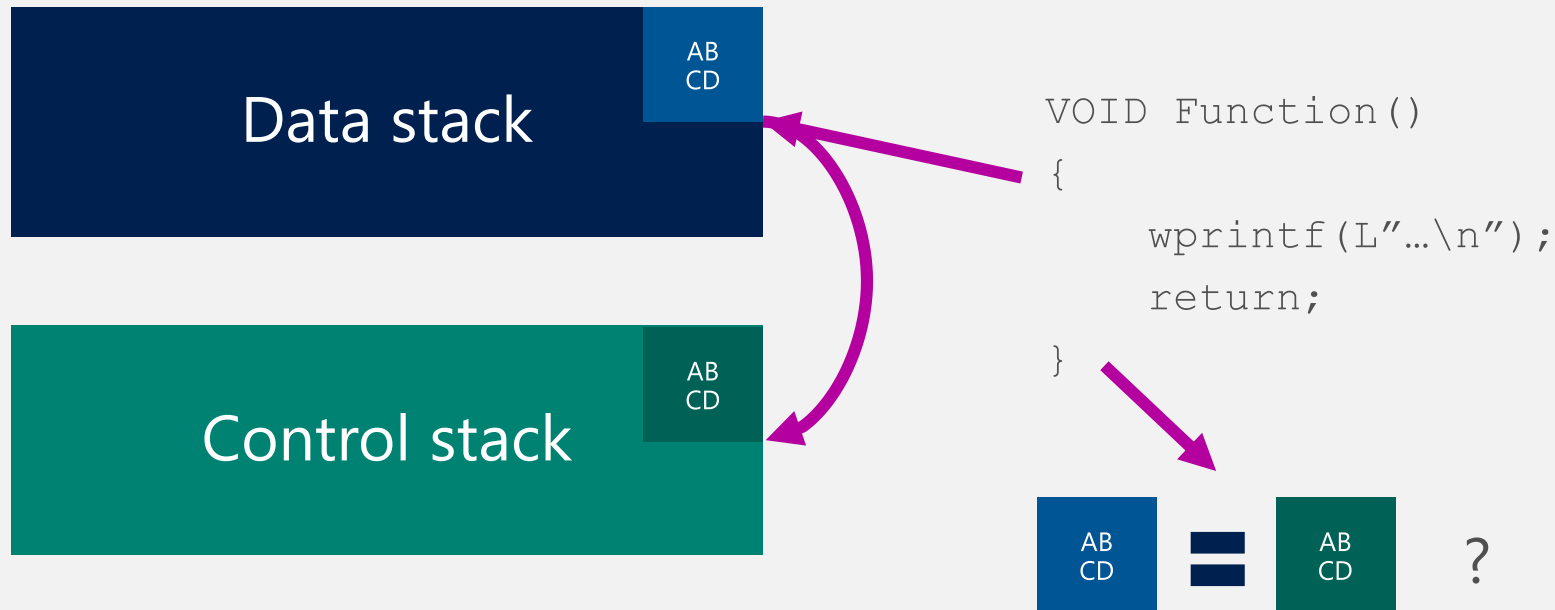
Preview specification:

<https://software.intel.com/sites/default/files/managed/4d/2a/control-flow-enforcement-technology-preview.pdf>

Return Flow Guard
(RFG)

Return Flow Guard (RFG)

RFG is an ABI-compliant shadow stack design that leverages x86_64 segmentation to hide the shadow stack

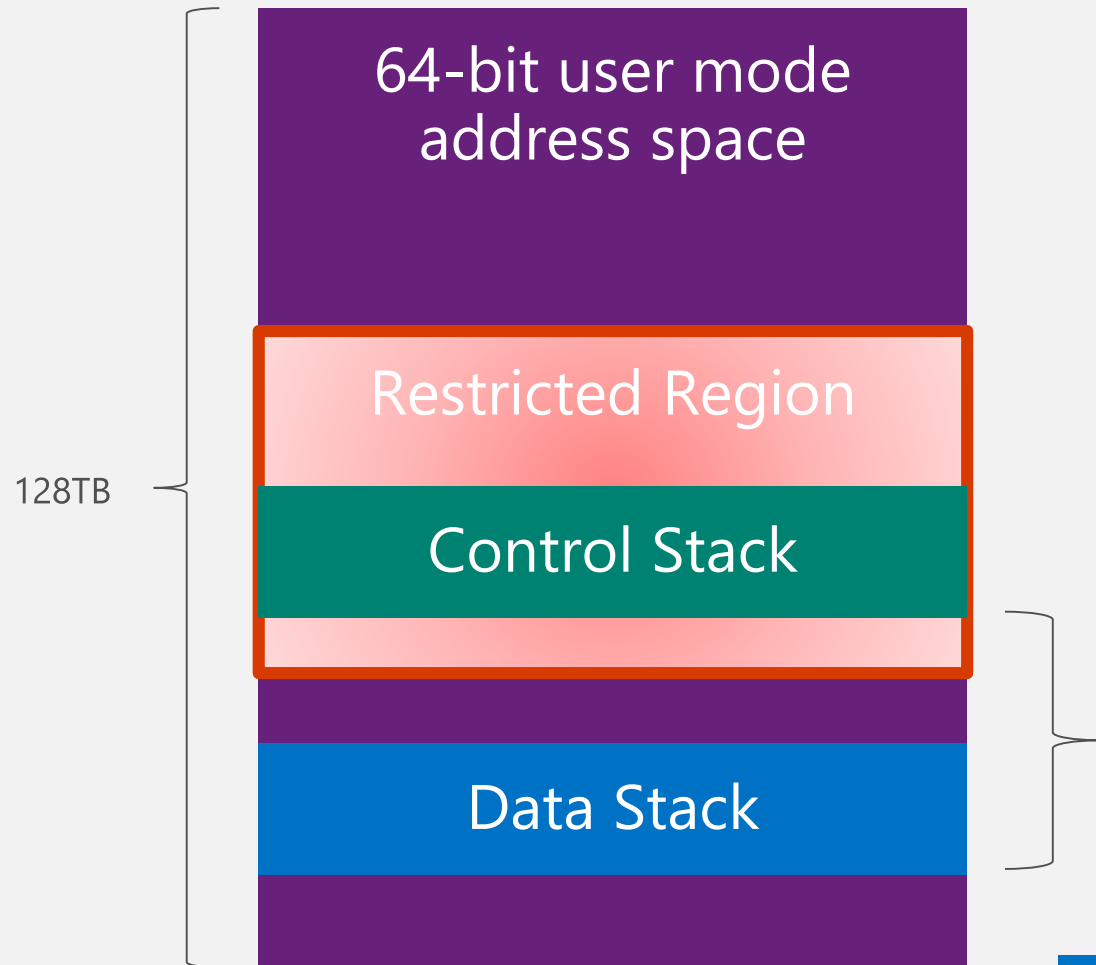


On function entry, a copy of the return address is stored in the control stack

On function exit, the return address in both stacks is compared

If they mismatch, terminate the process

RFG key design concepts



Per-function instrumentation is materialized at page-in

Prologue (9 bytes)

488b0424	mov	rax,qword ptr [rsp]
6448890424	mov	qword ptr fs:[rsp],rax

Epilogue (15 bytes)

644c8b1c24	mov	r11,qword ptr fs:[rsp]
4c3b1c24	cmp	r11,qword ptr [rsp]
0f8518a00300	jne	DLL!__guard_ss_verify_failure

FS segment base set to displacement between control & data stack

Thread->UserFsBase = ControlStack.BaseAddress - DataStack.BaseAddress

Context
switch

```
mov    eax, ThUserFsBase[rsi] ; load the bottom of the control stack delta FS base address
mov    edx, ThUserFsBase + 4[rsi] ; load the top of the control stack delta FS base address
mov    ecx, MSR_FS_BASE ; get the FS base MSR number
wrmsr  ; write the control stack delta FS base address
```

No pointers to a control stack exist in user mode memory

RFG appeared to meet all of our requirements

✖	Security	Must be robust against our threat model
✓	Performance	Cost must be within reason and proportional to value
✓	Compatibility	Don't break legacy apps
✓	Interoperability	Binary compatibility with previous <i>and</i> future versions of Windows
✓	ABI compliant	We can't rebuild the world
✓	Agility	Don't paint ourselves into a corner
✓	Developer friction	Cost for developers to enable should be minimal (ideally zero)

But there was just one big problem...

Our red team found a critical weakness

Using real vulnerabilities to evaluate RFG

Array.prototype.slice() Type Confusion

Takes a JavaScriptArray as input, returns a newly created JavaScriptArray

```
var fruits = ['Banana', 'Orange', 'Lemon', 'Apple', 'Mango'];  
var citrus = fruits.slice(1, 3);  
  
// fruits contains ['Banana', 'Orange', 'Lemon', 'Apple', 'Mango']  
// citrus contains ['Orange', 'Lemon']
```

MS16-084	Scripting Engine Memory Corruption Vulnerability	CVE-2016-3260	Jordan Rabet of Microsoft Offensive Security Research Team
----------	---	---------------	--

<https://github.com/Microsoft/ChakraCore/commit/17f3d4a4852dcc9e48de7091685b1862afb9f307>

Exploit Sequence

- Create a NativeFloatArray fake_obj_arr which contains a fake object
- Use CVE-2016-3260 to leak its address
- Use CVE-2016-3260 to create a JavaScript variable fake_obj pointing to fake_obj_arr + data_offset
- ...and a few other steps we won't be talking about today 😊

```
// build a fake dataview object
// could do this with a typedbuffer but would need to have the vtable address probably
fake_object_arr[0] = hex2float(0) // vtable
fake_object_arr[1] = hex2float(fake_type_addr) // type
fake_object_arr[2] = hex2float(0) // auxSlots
fake_object_arr[3] = hex2float(0) // objectArray / flags
fake_object_arr[4] = hex2float(0xffffffff) // length
fake_object_arr[5] = hex2float(locate_object(real_ab)) // arrayBuffer
fake_object_arr[6] = hex2float(0) // byteOffset
fake_object_arr[7] = hex2float(0xdeadbabedad) // buffer pointer
// create fake type for dataview
fake_object_arr[fake_type_offset] = hex2float(56) // TypeIds_DataView
fake_object_arr[fake_type_offset + 1] = hex2float(fake_object_start_addr) // javascriptLibrary : needs to be valid ptr

var fake_dv = create_offsetted_object(fake_object_arr, arr_inlineslots_offset)
```

Exploit was used as a platform to confirm a design-level bypass for RFG

Sticking to our threat model

The attack against RFG violated the assumptions of our threat model

- ✓ Accepting a weaker threat model would mean significant drop in ROI
- ✓ Long-term impact on cost to exploit would not be high enough to warrant long-term cost of feature
- ✓ Red team was key to vetting the solution
- ✓ @zer0mem also reported some interesting attacks 😊

RFG remains a research project; CET not affected

- ✓ Reminder: Microsoft has an ongoing \$100,000 USD bounty for defensive ideas 😊
- ✓ <https://technet.microsoft.com/en-us/security/dn425049.aspx>
- ✓ <https://blogs.technet.microsoft.com/srd/2015/09/08/what-makes-a-good-microsoft-defense-bounty-submission/>

Windows 10 1703 supports ACG, CIG, and CFG – RFG is not supported

Enabling code execution mitigations

Windows 10 provides multiple ways to enable mitigations for an app

Opt-in method	Details	Precedence
Runtime API	See <code>SetProcessMitigationPolicy</code>	4
EXE flag	Mitigation-specific flag in PE headers	3
Process creation attribute	See <code>UpdateProcThreadAttribute</code> with <code>PROC_THREAD_ATTRIBUTE_MITIGATION_POLICY</code>	2
Image File Execution Option (IFEO)	<code>Set MitigationOptions (REG_BINARY)</code> IFEO for an EXE Same bits as process creation attribute	1

Mitigation	Relevant policies	Runtime API	EXE bit	Process creation attribute	IFEO
CIG	<code>ProcessSignaturePolicy</code>	Yes	No	Yes	Yes
ACG	<code>ProcessDynamicCodePolicy</code>	Yes	No	Yes	Yes
CFG	<code>ProcessControlFlowGuardPolicy</code>	No	Yes (/guard:cf)	Yes	Yes
No Child Process	<code>CHILD_PROCESS_POLICY</code> process creation attribute	No	No	Yes	No

See also: <https://docs.microsoft.com/en-us/windows/threat-protection/overview-of-threat-mitigations-in-windows-10>

Checking if mitigations are enabled



Handy PowerShell cmdlets can be used to see what mitigations are enabled

```
Administrator: Windows PowerShell  
Windows PowerShell  
Copyright (C) 2016 Microsoft Corporation. All rights reserved.  
PS C:\WINDOWS\system32> Install-Module -Name ProcessMitigations
```



```
Administrator: Windows PowerShell  
PS C:\WINDOWS\system32> Get-ProcessMitigation -Name vmwp.exe -RunningProcess  
  
ProcessName: vmwp  
Source      : Running Process  
Id          : 2836  
  
DEP:  
  Enable           : ON  
  Disable ATL      : OFF  
  
ASLR:  
  BottomUp         : ON  
  ForceRelocate    : ON  
  HighEntropy      : ON  
  DisallowStripped : OFF  
  
StrictHandle:  
  RaiseExceptionOnInvalid : ON  
  HandleExceptionsPermanently : ON  
  
System Call:  
  DisallowWin32kSysCalls : ON  
  
ExtensionPoint:  
  DisableExtensionPoints : ON  
  
DynamicCode:  
  ProhibitDynamicCode : ON  
  AllowThreadOptOut   : OFF  
  AllowRemoteDowngrade : OFF  
  
CFG:  
  EnableCFG : ON  
  EnableExportSuppression : OFF  
  StrictMode : ON  
  
BinarySignature:  
  MicrosoftSignedOnly : ON  
  StoreSignedOnly : OFF  
  MitigationOptIn : OFF  
  
FontDisable:  
  DisableNonSystemFonts : ON  
  AuditNonSystemFontLoading : OFF  
  
ImageLoad:  
  NoRemoteImages : ON  
  NoLowMandatoryLabelImages : ON  
  PreferSystem32Images : ON
```



Protected apps in Windows 10 1703

	 Microsoft Edge	 Hyper-V VMWP
ACG	✓	✓
CIG	✓	✓
CFG	✓	✓
No Child Process	✓	✓

- For Microsoft Edge, ACG and CIG are disabled in the presence of certain 3rd party extensions (e.g. IMEs, legacy graphics drivers). ACG is currently enabled only on 64-bit.
- For Hyper-V VMWP, the no child process mitigation is disabled when certain features are used (vTPM)

What about the Windows kernel?

Windows 10 leverages Hyper-V Virtualization Based Security (VBS) to enable CI, ACG, and CFG for the Windows kernel

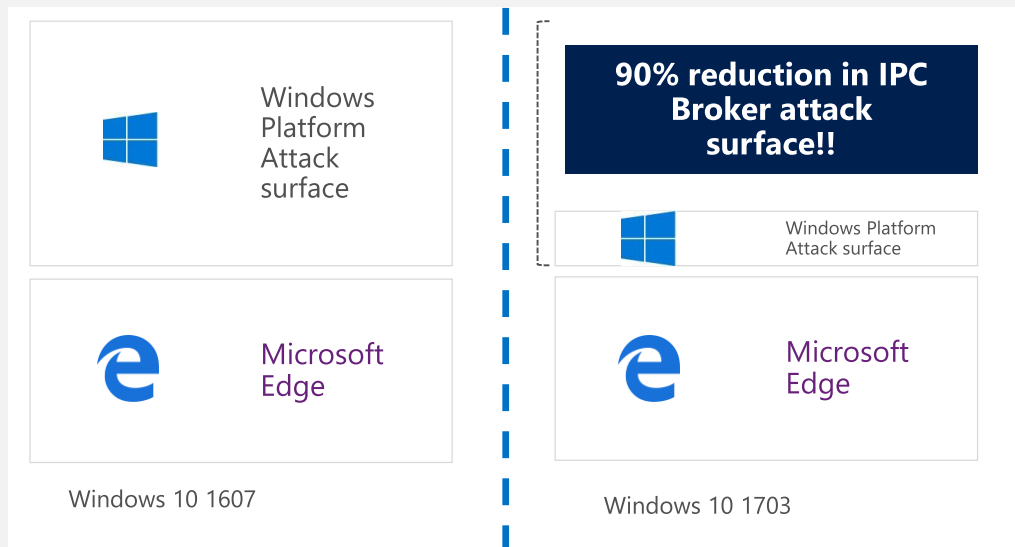
Kernel Mode		 Windows 10 1703
	ACG	 Enabled with Hypervisor-Enforced Code Integrity (HVCI) via Device Guard
	CIG	
	CFG	

Conclusion & next steps

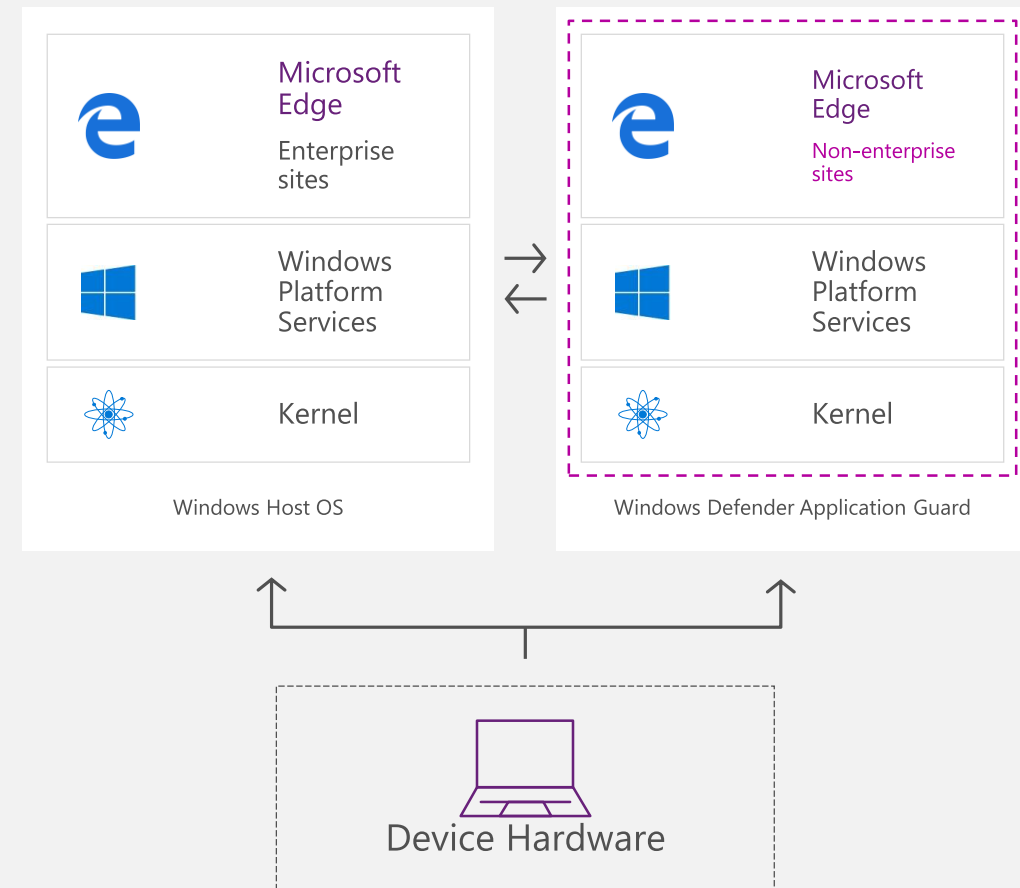
It's not RCE mitigations *or* sandboxing – it's both

We are continuing to invest in improved isolation technologies as part of our layered strategy

Improved software isolation (Microsoft Edge AppContainer Profile)



Virtualized Isolation (Application Guard)



To learn more:

<https://blogs.windows.com/msedgedev/2017/03/23/strengthening-microsoft-edge-sandbox>

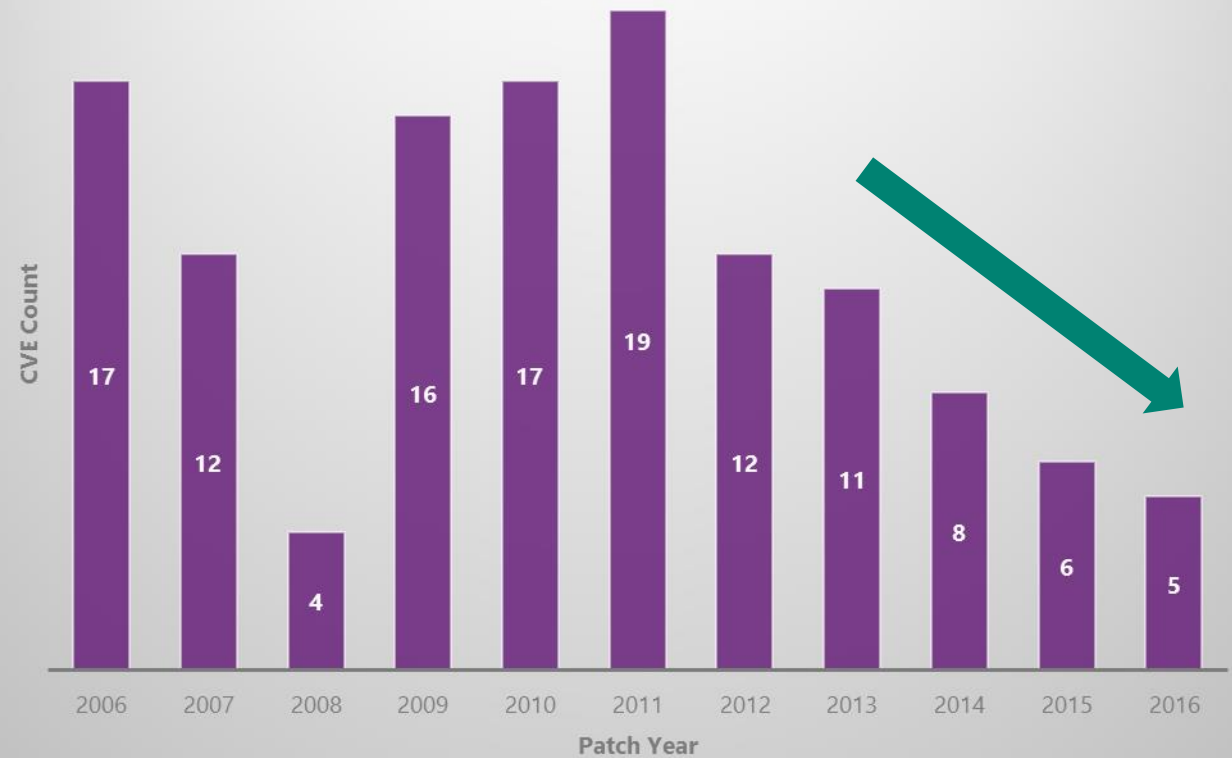
Our work is having a measurable impact

of known zero day RCE exploits has declined despite increase in known vulnerabilities

of Microsoft Remote Code Execution CVEs by year patched



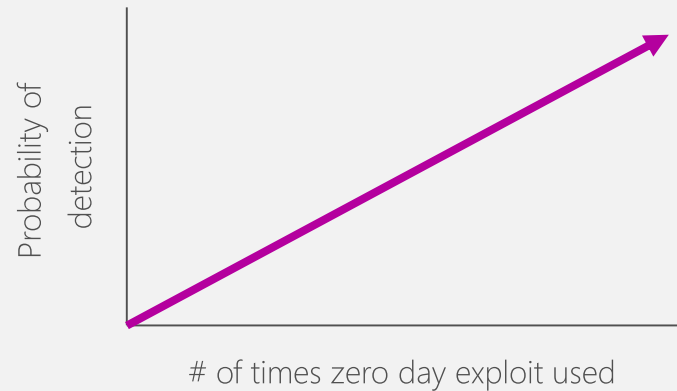
of Microsoft Remote Code Execution CVEs exploited as zero day by year patched



What about the zero days we don't know about?

We cannot effectively measure the number of unknown zero day exploits, but we don't really need to

Hypothesis: increased exploit costs drive selective use



- ✓ Probability of detection increases with zero day use
 - Attackers are incentivized to minimize use
 - Targets that detect zero day may alert vendor
- ✓ Selective use reduces downstream supply
 - Many actors lack means and capability to acquire

Assertion: economics of the zero day market have shifted

- ✓ Windows 10 is always up to date
 - Poor ROI for exploiting patched vulnerabilities
 - Rapid evolution of defensive technologies
- ✓ Mass-market exploit kits struggling to maintain supply
 - Decrease in reusable & public exploits
 - Cost to acquire exceeds expected ROI
- ✓ Market has shifted toward social engineering
 - Macros, tech support scams, so on

MAGNITUDE ACTOR ADDS A SOCIAL ENGINEERING SCHEME FOR WINDOWS 10

MARCH 08, 2017 Kafeine <https://www.proofpoint.com/us/threat-insight/post/magnitude-actor-social-engineering-scheme-windows-10>

Rapid evolution of defense in Windows 10

New & improved defenses are shipping with nearly every Windows Insider Preview (WIP) build



Matt Miller
@epakskape



vmwp.exe (Hyper-V VM worker proc) enables nearly all mitigations in Windows 10 1703. Easy to see using this cmdlet: powershellgallery.com/packages/Proce...

1:21 PM - 10 Apr 2017



Dave dwizzle Weston
@dwizzleMSFT



If you instal Windows 10 1703 and enable HVCI we now also enable kCFG. Kernel mode exploits are subject to control flow integrity in Win10

7:57 PM - 14 Apr 2017

62 65



Dave dwizzle Weston
@dwizzleMSFT



Hypervisor Isolation of Edge is available in this weeks flight!!
Grab it, play with it, report bugs.

#WindowsDefenderApplicationGuard #WDAG
twitter.com/windowsinsider...

10:08 AM - 4 May 2017

76 76



Matt Miller
@epakskape



Say goodbye to exploit techniques that rely on the fixed HAL heap kernel mapping with the Creators Update of Windows 10 64-bit twitter.com/bluefrostsec/s...

7:16 AM - 11 May 2017

57 69

Live life in the WIP
Fast Ring to get the
latest defenses 😊

Windows 10 build 16179 has a subtle heap randomization improvement that breaks this primitive:

github.com/saaramar/Deter...

3:13 PM - 26 Apr 2017

86 109

Lessons learned

3rd party binary extensions are security's nemesis

- Very difficult to maintain compatibility and interoperability with 3rd party binary extensions
- Apps need to tightly control their binary extension points to maintain security agility

We're not done yet with mitigating arbitrary native code execution

- CET should provide robust protection for return addresses once available

Layered, data-driven, and red-team-assisted defense is key

- No silver bullets exist in security
- Focus investments on defensible positions

What's next?

We believe our strategy is working & are continuing to execute on it

Kill more bug
classes

Drive down
attack surface

Break more
exploitation
techniques

Improve least
privilege
containment

Lots of exciting things in store. Stay tuned to the WIP Fast Ring!

Fascinated by what you saw? Want to help us make the online world safer?



aka.ms/bugbounty

Report vulnerabilities &
mitigation bypasses via our
bounty programs!

<https://aka.ms/bugbounty>

Or come work with us. We're hiring 😊

<https://aka.ms/cesecurityopenjobs>

<https://aka.ms/wdgsecurityjobs>

Bounty program update!

Hyper-V & Mitigation Bypass Bounty Updates

We've clarified and expanded the scope of our Hyper-V bounty program

- Increased Hyper-V payout up to \$150,000 USD
- Denial of Service and Information Disclosure are now in scope for the Hyper-V bounty
- Clarified the payout structure for the Hyper-V and Mitigation Bypass bounty programs
- Hyper-V bounty is now separate from the Mitigation Bypass Bounty



Hyper-V bounty page: <https://technet.microsoft.com/en-us/security/mt784431>

Mitigation bypass bounty page: <https://technet.microsoft.com/en-us/security/dn425049>

Acknowledgements

Tons of people worked hard to help make this possible!

Thanks to the following groups & teams for their contributions:

- Windows kernel and Hyper-V teams
- Visual C++ team
- Microsoft Edge security team
- WDG Offensive Security Research team
- Microsoft Security Response Center
- And everyone else who contributed!

